
[SoC Design – Term Project]

Accelerator for CNN

Chester Sungchung Park
SoC Design Lab, Konkuk University
Webpage: <http://soclab.konkuk.ac.kr>

Outline

- ☐ Hardware accelerator
- ☐ Convolution in CNN
 - AlexNet for ImageNet
- ☐ Design flow
- ☐ Design constraints
- ☐ Evaluation
- ☐ Submission
- ☐ Presentation
- ☐ Appendix

Convolutional Neural Networks (CNN)

□ Example: AlexNet for ImageNet

Sample Image



Classifications

Football helmet

Broccoli (973)


Suit, Suit of Clothes

Baseball

ImageNet (Input Images)




Convolution in CNN



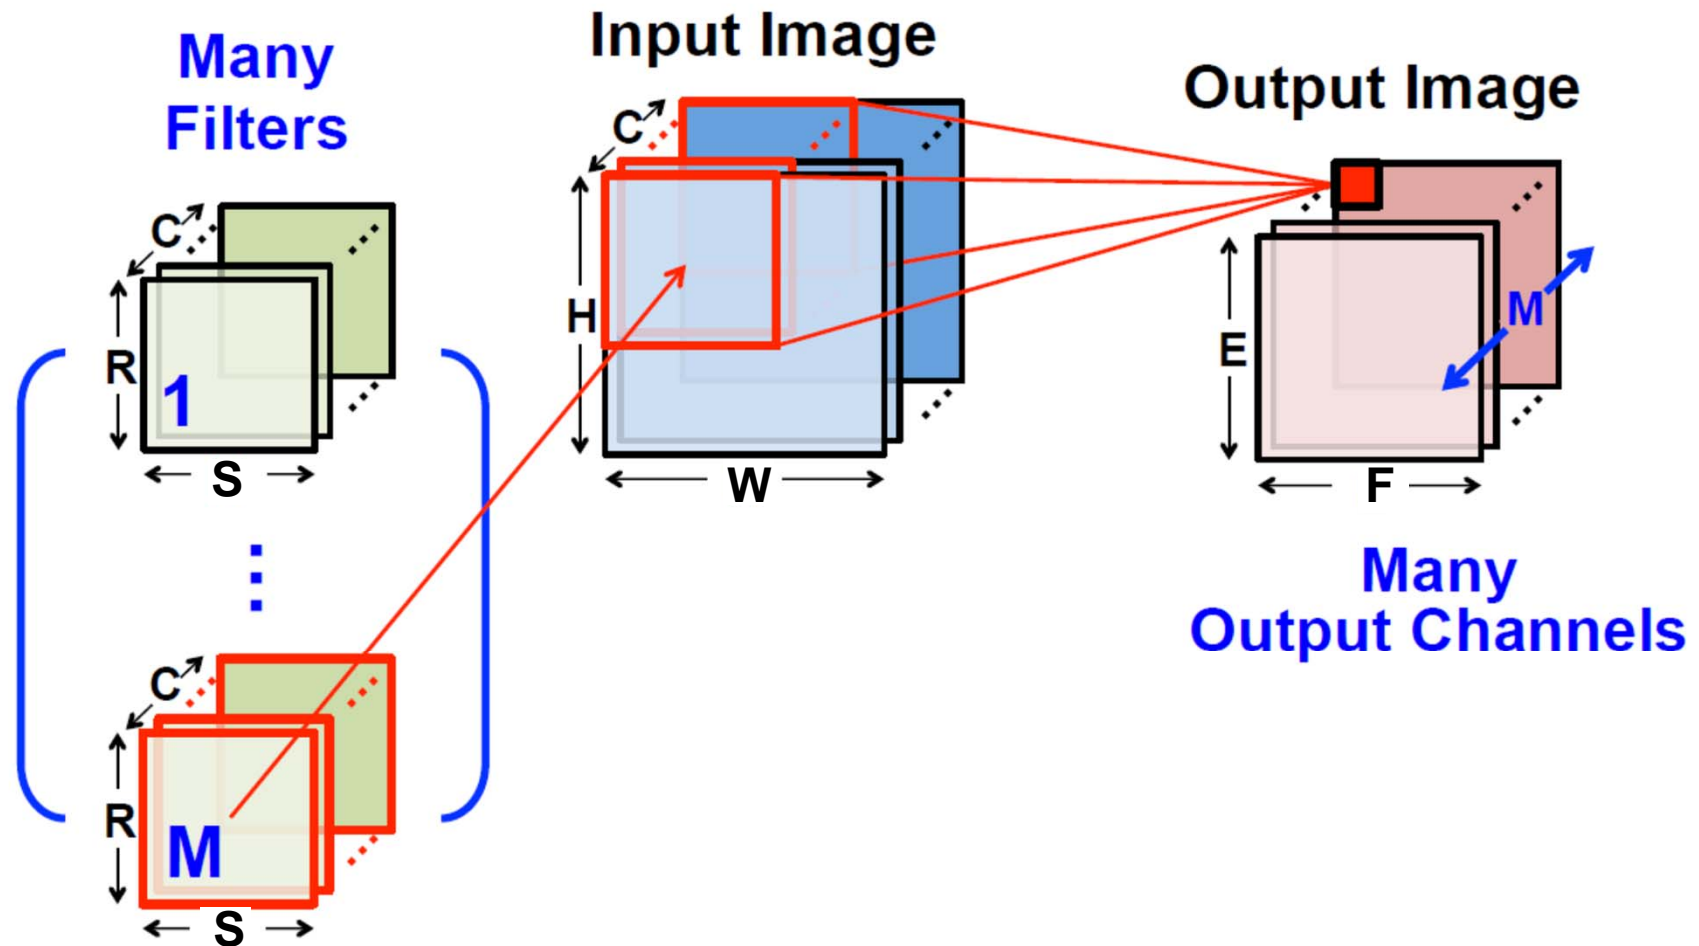
Source: Y-H. Chen

Convolution in CNN



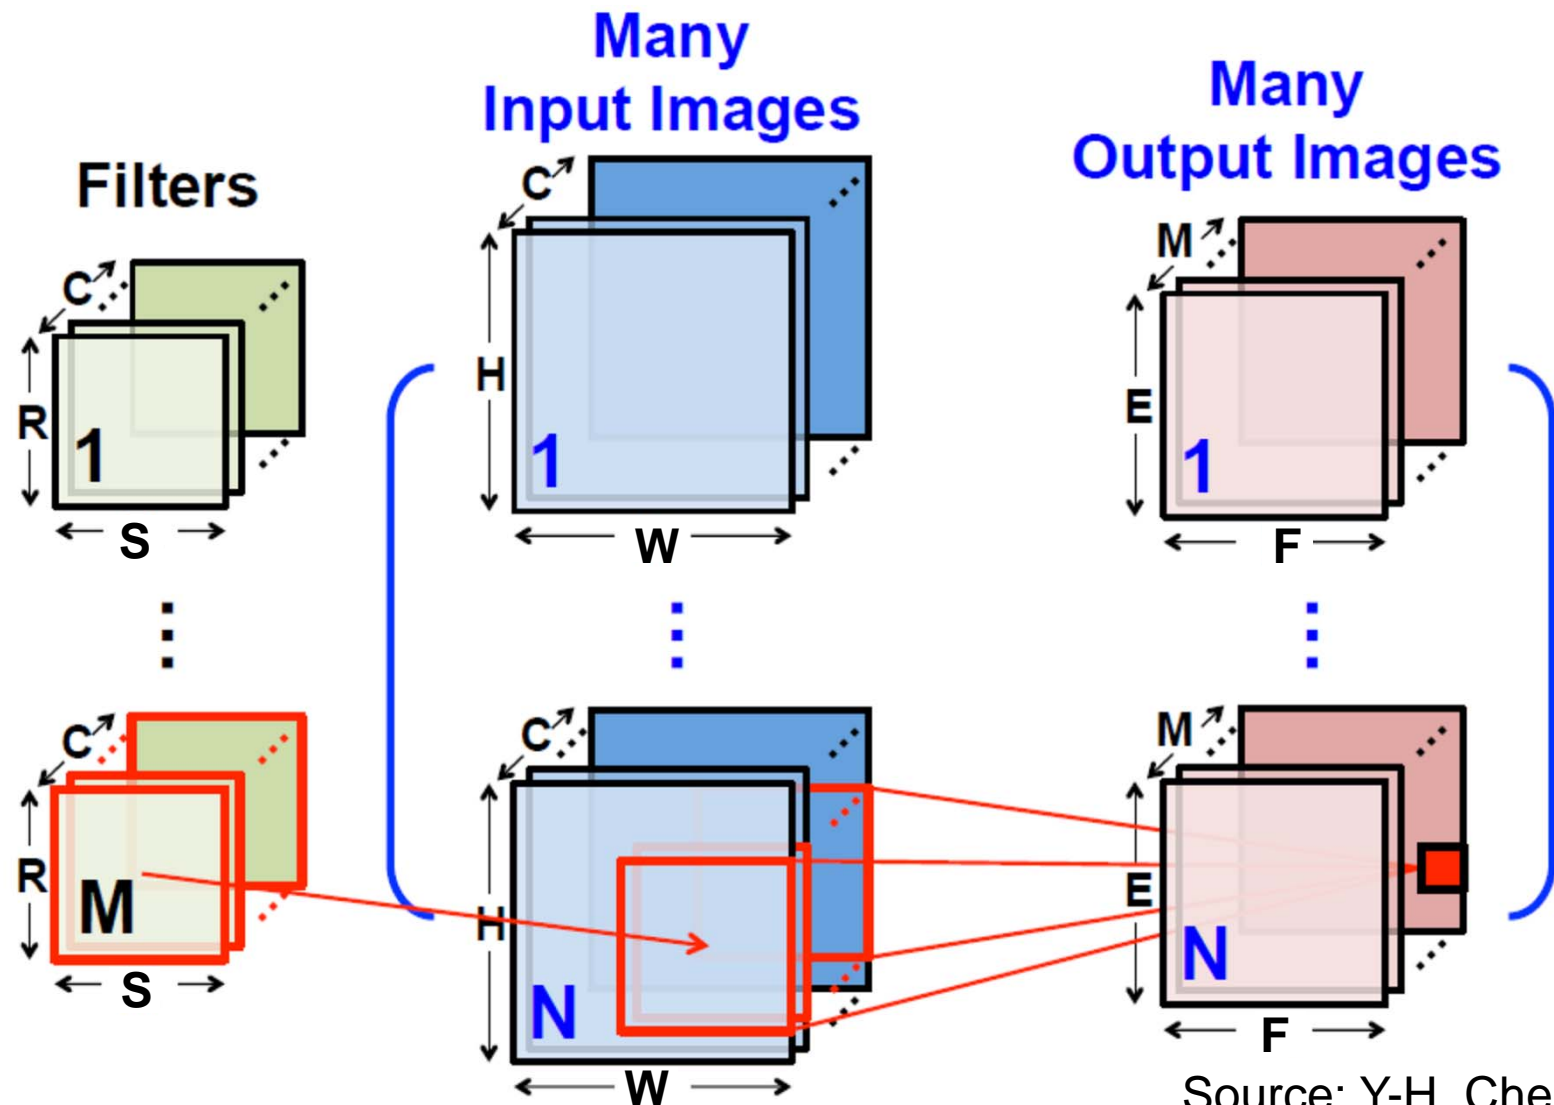
Source: Y-H. Chen

Convolution in CNN



Source: Y-H. Chen

Convolution in CNN



Source: Y-H. Chen

Convolution in CNN


□ Pseudo code

```
For (n = 0; n < N; n++) → Batch
  For (c = 0; c < C; c++) → Channel
    For (m = 0; m < M; m++) → Filter
      For (f = 0; f < F; f++) → of height
        For (e = 0; e < E; e++) → of width
          For (s = 0; s < S; s++) → f height
            For (r = 0; r < R; r++) → f width
              of[e][f][m][n] += if[r + e][s + f][c][n] · f[r][s][c][m]
```

Feature map out Feature map in Filter

Convolution in CNN

□ AlexNet for ImageNet



Convolution in CNN

□ AlexNet for ImageNet (cont'd)

		Width (W/S)	Height (H/R)	# I-Ch (C)	# O-Ch (M)
Input Image		227	227	3	-
Filter	Layer 1	11		3	96
	Layer 2	5		48	256
	Layer 3	3	3	256	384
	Layer 4	3	3	192	384
	Layer 5	3	3	192	256
	Layer 6	6	6	256	4096
	Layer 7	1	1	4096	4096
	Layer 8	1	1	4096	1000

Reference Implementation


- ❑ Run the attached reference implementation as explained in the appendix
 - AlexNet with layer 3 in HW and the others in SW
 - No local reuse (NLR) chosen as dataflow


source_TP_

- ❑ For details of NLR, refer to the following paper
 - “Optimizing FPGA-based Accelerator Design for Deep Convolutional Neural”, C. Zhang et al., FPGA 2015

No Local Reuse (NLR)

❑ Original algorithm



```

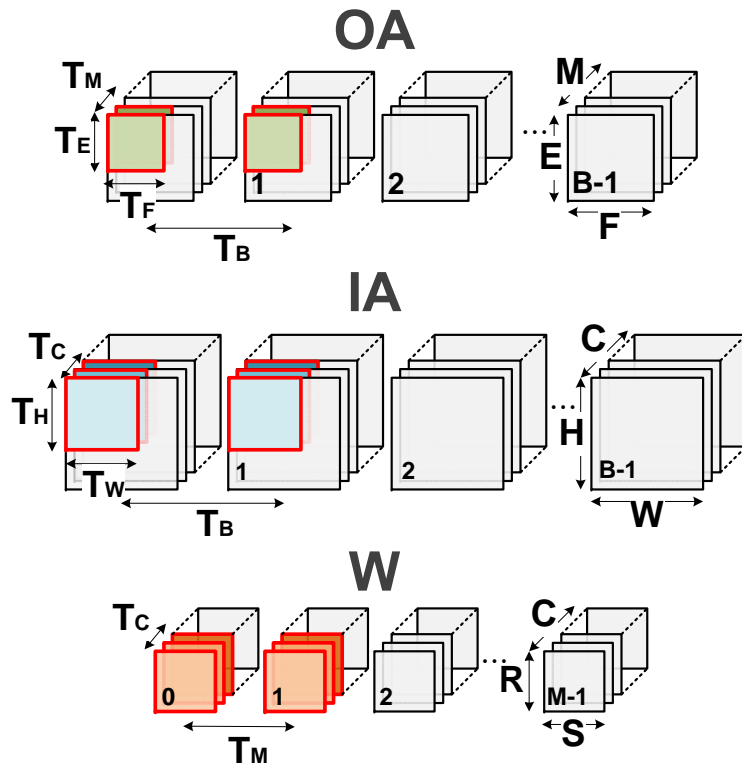
for (b = 0; b < B; b++) {
  for (e = 0; e < E; e++) {
    for (f = 0; f < F; f++) {
      for (m = 0; m < M; m++) {
        for (c = 0; c < C; c++) {
          for (r = 0; r < R; r++) {
            for (s = 0; s < S; s++) {
              O[b][m][e][f] +=
                W[m][c][r][s] * I[b][c][e+r][f+s];
            } } } } } } }

```

B: No. of images
 E: Output height
 F: Output width
 M: No. of filters
 C: No. of features
 R: Filter height
 S: Filter width

No Local Reuse (NLR)

□ Loop tiling for NLR



```

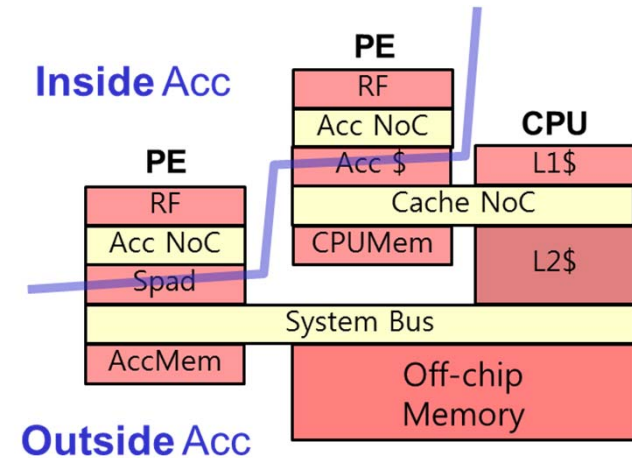
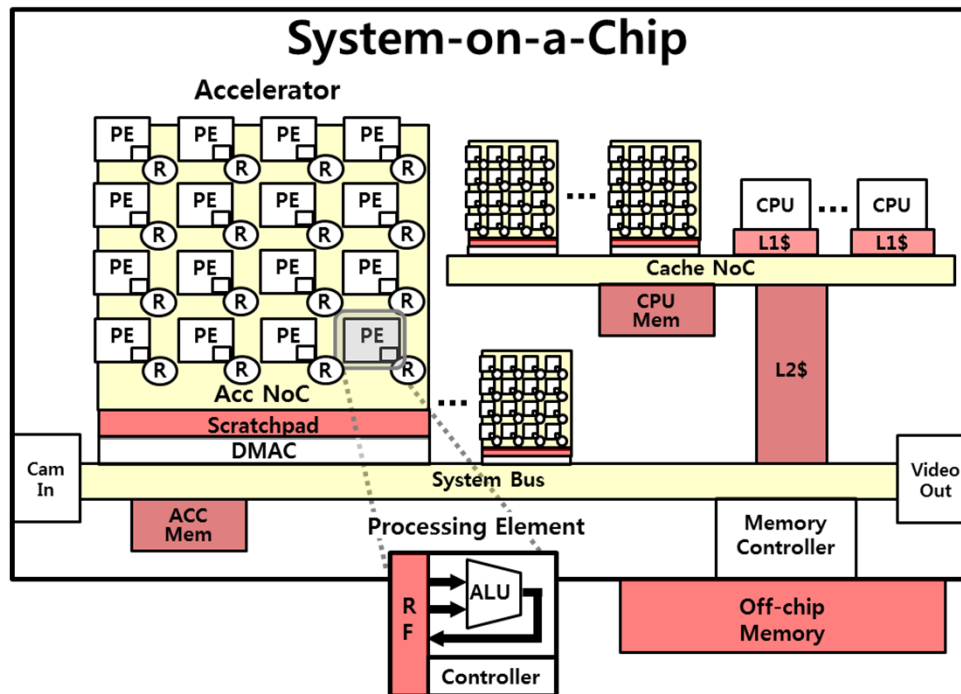
for (b = 0; b < B; b += TB) {
  for (e = 0; e < E; e += TE) {
    for (f = 0; f < F; f += TF) {
      for (m = 0; m < M; m += TM) {
        for (c = 0; c < C; c += TC) {
          Ibuf = I[b:b+TB-1][c:c+TC-1][e:e+TE+R-2][f:f+TF+S-2];
          Wbuf = W[m:m+TM-1][c:c+TC-1][0:R-1][0:S-1];
          for (r = 0; r < R; r++) {
            for (s = 0; s < S; s++) {
              for (tb = 0; tb+b < min(B,b+TB); tb++) {
                for (te = 0; te+e < min(E,e+TE); te++) {
                  for (tf = 0; tf+f < min(F,f+TF); tf++) {
                    for (tm = 0; tm < TM; tm++) {
                      for (tc = 0; tc < TC; tc++) {
                        wx[tm][tc] = Wbuf[tm][tc][r][s];
                        ix[tc] = Ibuf[tb][tc][te+r][tf+s];
                        ox[tm] = Obuf[tb][tm][te][tf];
                        Obuf[tb][tm][te][tf] = wx[tm][tc]*ix[tc]+ox[tm];
                      }
                    }
                  }
                }
              }
            }
          }
          O[b:b+TB-1][m:m+TM-1][e:e+TE-1][f:f+TF-1] = Obuf;
        }
      }
    }
  }
}

```

TB: Batch size
 TE: Output height /tile
 TF: Output width /tile
 TM: No.of filters /tile
 C: No.of features /tile
 R: Filter height /tile
 S: Filter width /tile

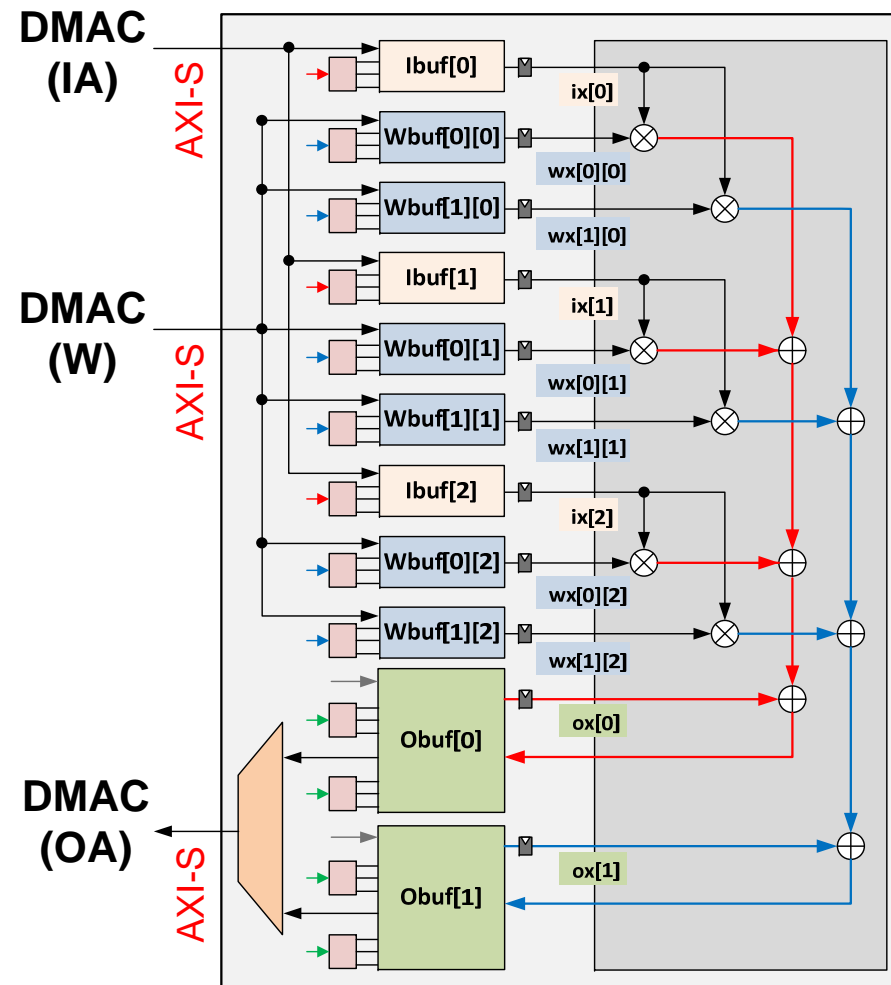
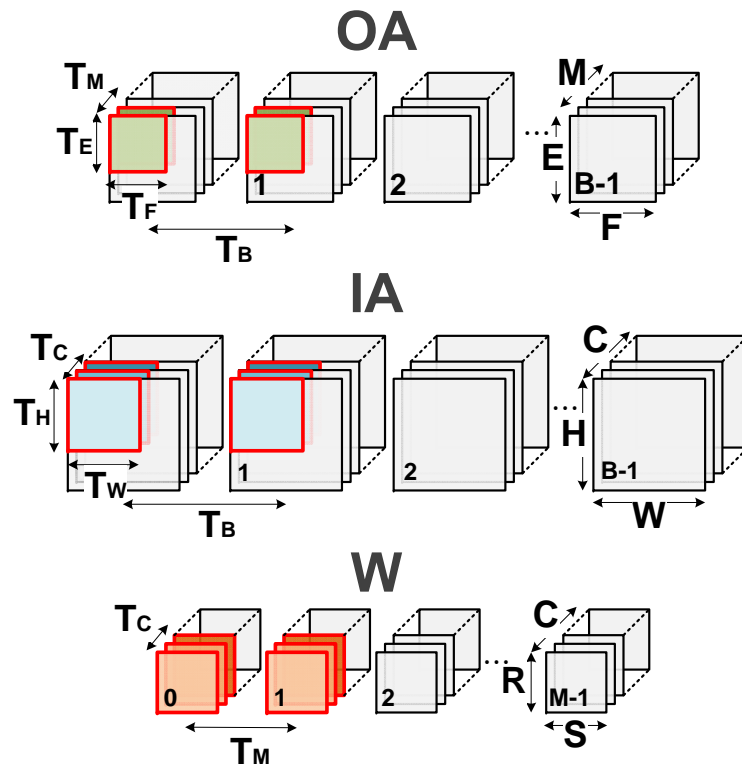
No Local Reuse (NLR)

❑ Memory hierarchy



No Local Reuse (NLR)

Accelerator hardware





* $T_C = 2, T_M = 3$

No Local Reuse (NLR)

```

for (b = 0; b < B; b += TB) {
  for (e = 0; e < E; e += TE) {
    for (f = 0; f < F; f += TF) {
      for (m = 0; m < M; m += TM) {
        for (c = 0; c < C; c += TC) {
          Ibuf = I[b:b+TB-1][c:c+TC-1][e:e+TE+R-2][f:f+TF+S-2];
          Wbuf = W[m:m+TM-1][c:c+TC-1][0:R-1][0:S-1];
          .....
        }
      }
    }
  }
  O[b:b+TB-1][m:m+TM-1][e:e+TE-1][f:f+TF-1] = Obuf;
}
    
```

**Outer-Pipeline
(Inter-Pass)**



* $C = 6, M = 6, T_C = 2, T_M = 3$


No Local Reuse (NLR)

```


for (b = 0; b < B; b += TB) {
  for (e = 0; e < E; e += TE) {
    for (f = 0; f < F; f += TF) {
      for (m = 0; m < M; m += TM) {
        for (c = 0; c < C; c += TC) {
          Ibuf = I[b:b+TB-1][c:c+TC-1][e:e+TE+R-2][f:f+TF+S-2];
          Wbuf = W[m:m+TM-1][c:c+TC-1][0:R-1][0:S-1];
          for (r = 0; r < R; r++) {
            for (s = 0; s < S; s++) {
              for (tb = 0; tb+b < min(B,b+TB); tb++) {
                for (te = 0; te+e < min(E,e+TE); te++) {
                  for (tf = 0; tf+f < min(F,f+TF); tf++) {
                    for (tm = 0; tm < TM; tm++) {
                      for (tc = 0; tc < TC; tc++) {
                        wx[tm][tc] = Wbuf[tm][tc][r][s];
                        ix[tc] = Ibuf[tb][tc][te+r][tf+s];
                        ox[tm] = Obuf[tb][tm][te][tf];
                        Obuf[tb][tm][te][tf] = wx[tm][tc]*ix[tc]+ox[tm];
                      }
                    }
                  }
                }
              }
            }
          }
          O[b:b+TB-1][m:m+TM-1][e:e+TE-1][f:f+TF-1] = Obuf;
        }
      }
    }
  }
}

```


**Outer-Pipeline
(Inter-Pass)**



**Inner-Pipeline
(Intra-Pass)**



Unroll



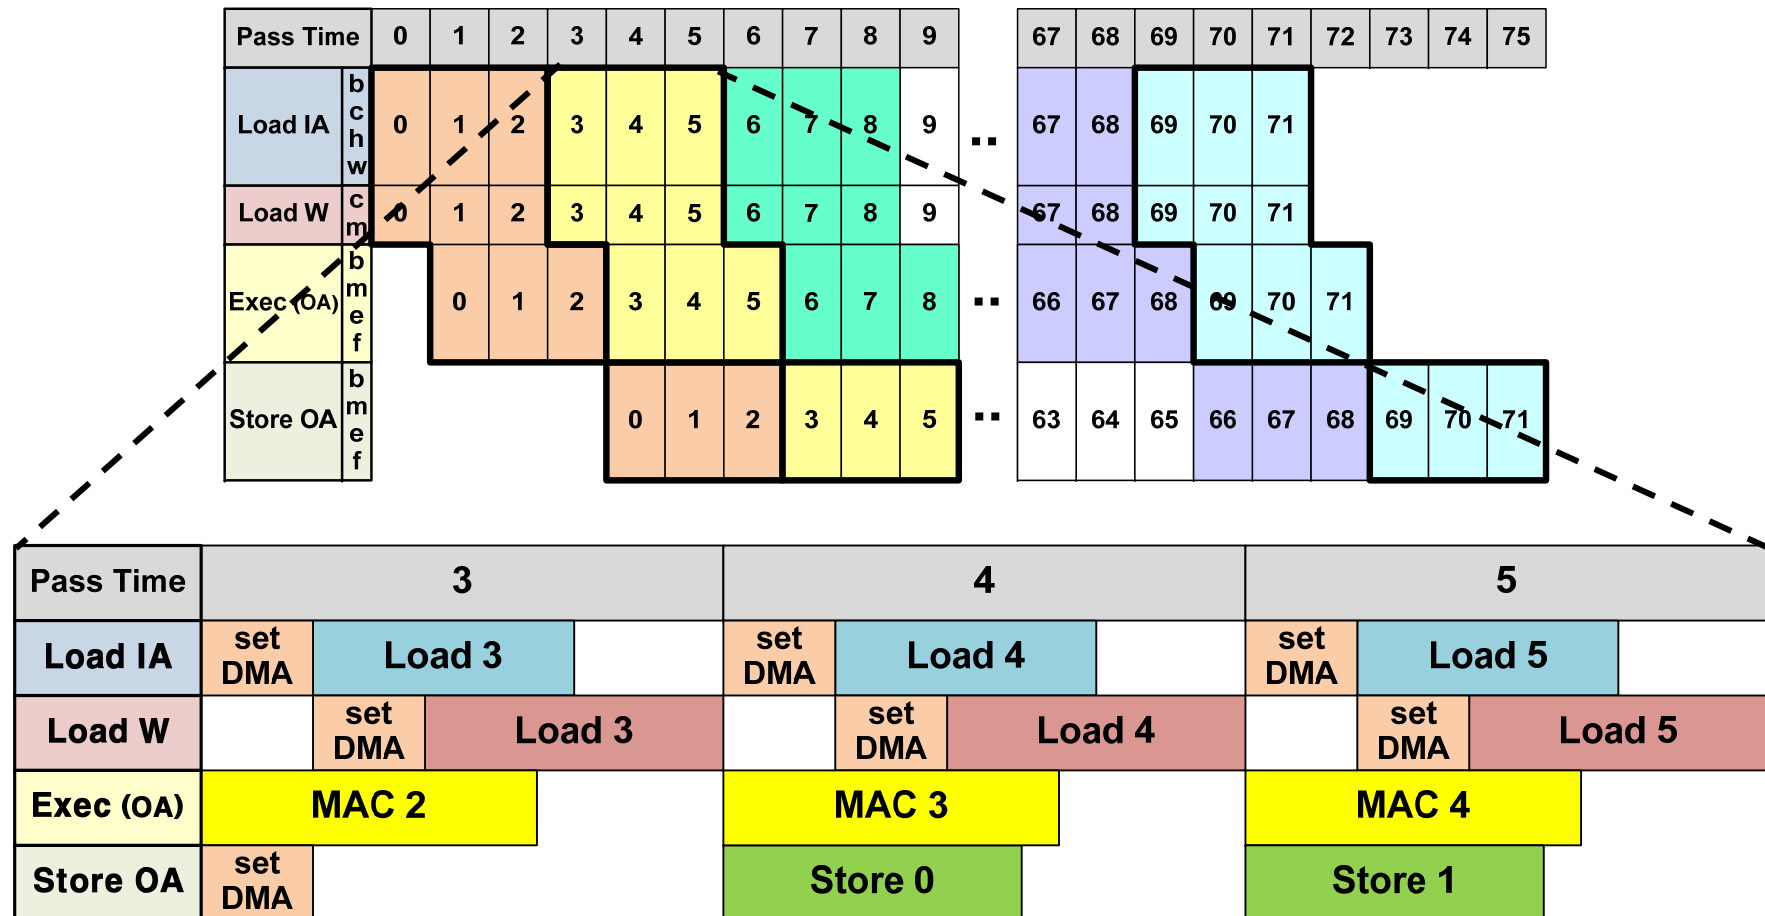
No Local Reuse (NLR)

❑ Outer pipeline (inter-pass)

Pass Time		0	1	2	3	4	5	6	7	8	9			67	68	69	70	71	72	73	74	75
Load IA	b	0	0	0	0	0	0	0	0	0	0	..		2	2	2	2	2				
	c	0	3	6	0	3	6	0	3	6	0			3	6	0	3	6				
	h	0	0	0	0	0	0	0	0	0	0			3	3	3	3	3				
	w	0	0	0	0	0	0	0	0	0	3			3	3	3	3	3				
Load W	c	0	3	6	0	3	6	0	3	6	0	..		3	6	0	3	6				
	m	0	0	0	2	2	2	4	4	4	0			2	2	4	4	4				
Exec (OA)	b			0		0			0			..		2				2				
	m			0		2			4					2				4				
	e			0		0			0					3				3				
	f			0		0			0					3				3				
Store OA	b					0			0			..		2				2			2	
	m					0			2					0				2			4	
	e					0			0					3				3			3	
	f					0			0					3				3			3	

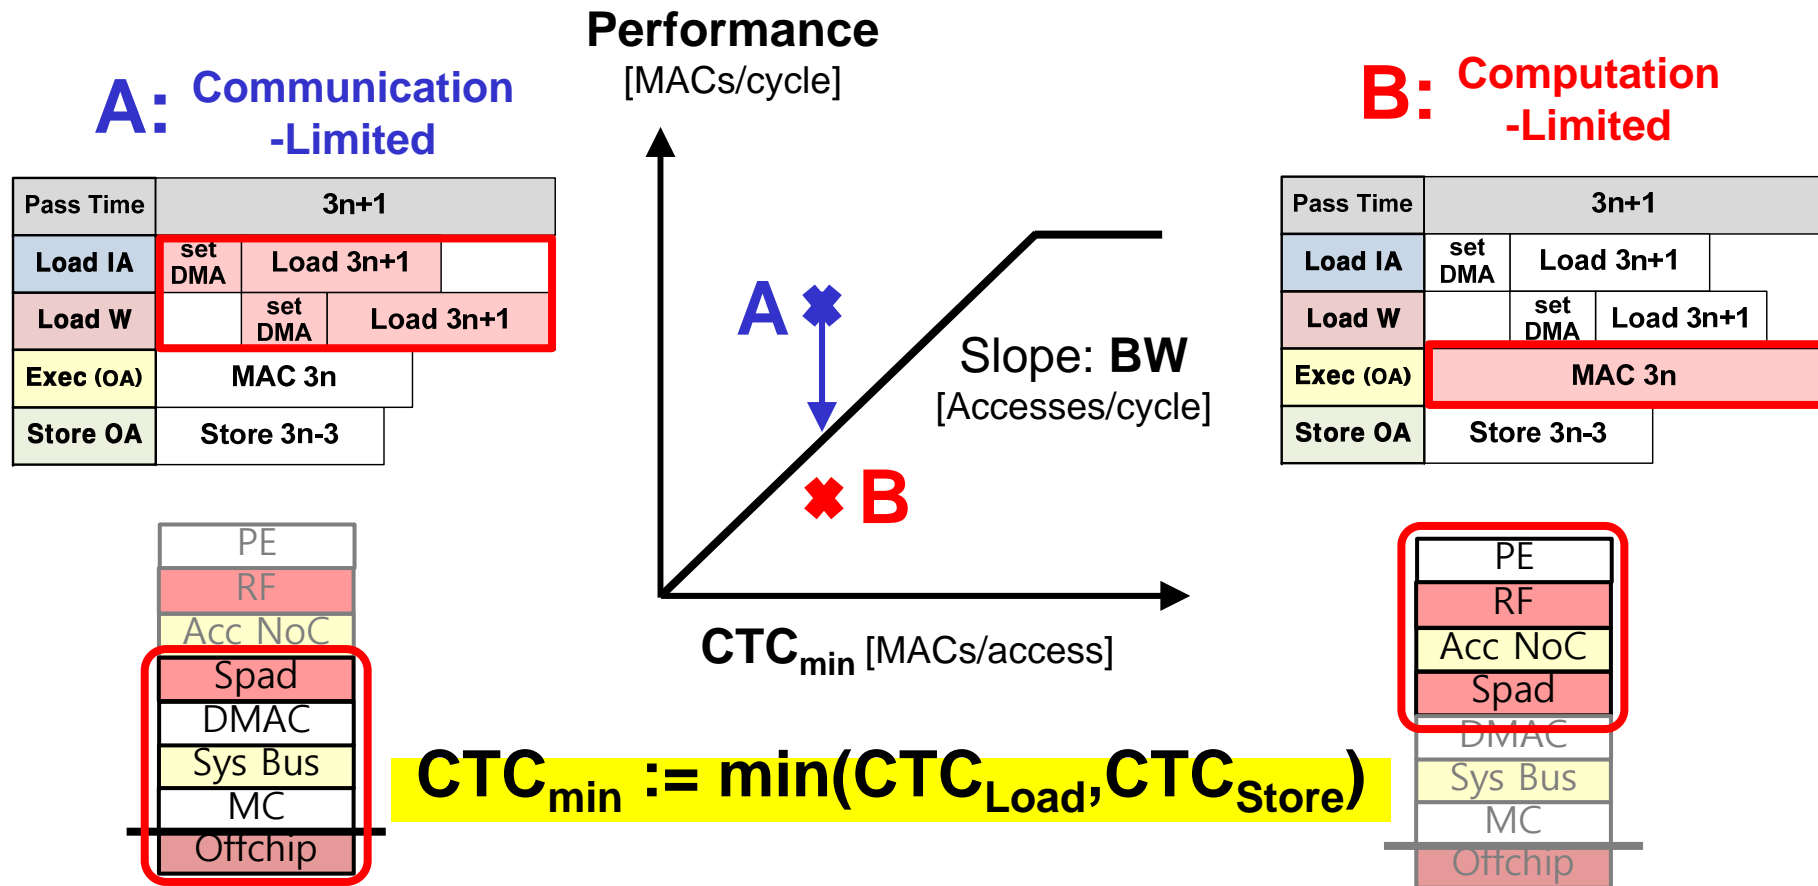
No Local Reuse (NLR)

❑ Outer pipeline (inter-pass) (cont'd)

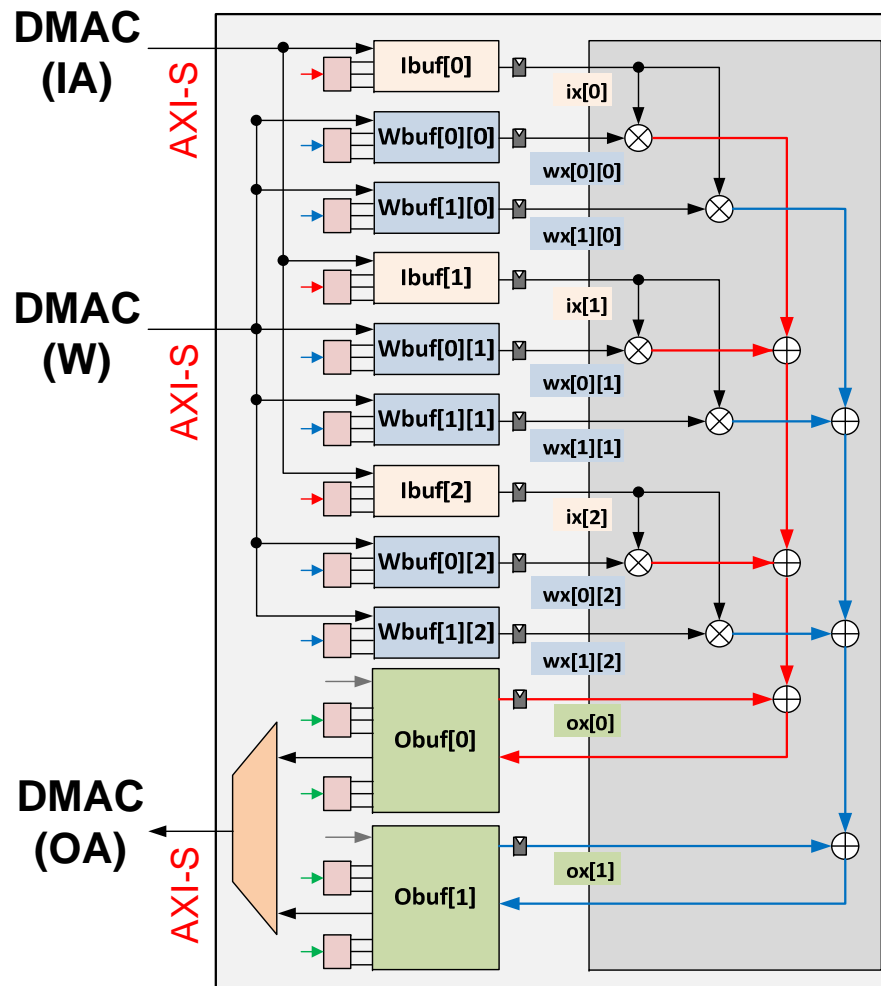


No Local Reuse (NLR)

Roofline model (revisited)



Available Hardware Blocks



Denoted as
"hls23x" for layer x

* $T_C = 2$, $T_M = 3$

Available Hardware Blocks


You're allowed to use **whatever hardware blocks**
(even to use **all** the blocks)
to improve the performance

$T_C \times T_M$	3 x 4	2 x 8	8 x 2	2 x 4	4 x 2
L1	hls341	NA	NA	NA	NA
L2	NA	hls282	hls822	hls242	hls422
L3	NA	hls283	hls823	hls243	hls423
L4	NA	hls284	hls824	hls244	hls424
L5	NA	hls285	hls825	hls245	Hls425

Will be available in the course webpage soon

Design Flow


Vivado



SDK



ZYNQ/ZedBoard



Design Constraints

- ❑ You are allowed to modify **only** the following function where **all** the multiply-accumulate (MAC) operations are implemented in **HW**

- **conv_hw**

- ❑ You are **not** allowed to change anything outside the above function


```
*main.c
113
114 void cnn_opt(short *ofmap, short *ifmap, int data_set)
115 {
116     // You are allowed to modify only the inside of this function
117     // From this line
118
119     int i = 0;
120
121
122     short *ofmap1 = (short *) calloc(E_C1*F_C1*M_C1*N_C1, sizeof(short));
123     short *ofmap2 = (short *) calloc(E_P1*F_P1*C_P1*N_P1, sizeof(short));
124     short *ofmap3 = (short *) calloc(E_C2*F_C2*M_C2*N_C2, sizeof(short));
125     short *ofmap4 = (short *) calloc(E_P2*F_P2*C_P2*N_P2, sizeof(short));
126     short *ofmap5 = (short *) calloc(E_C3*F_C3*M_C3*N_C3, sizeof(short));
127     short *ofmap6 = (short *) calloc(E_R1*F_R1*M_R1*N_R1, sizeof(short));
128
129     //Input mapping
130     for (i = 0; i < H_C1*W_C1; i++) {
131         ifmap[i] = (short)(data[data_set][i] * pow(2, SCALE - 1));
132     }
133
134     // Layer #1
135     //
136     convolution(ofmap1, ifmap, fmap1, N_C1, C_C1, M_C1, F_C1, E_C1, R_C1, S_C1, H_C1, W_C1, U_C1);
137     bias(ofmap1, ofmap1, bias1, N_C1, M_C1, E_C1, F_C1);
138     pool(ofmap1, ofmap1, E_C1, M_C1, C_C1);
139
140     // Layer #2
141     //
142     convolution(ofmap3, ofmap2, fmap2, N_C2, C_C2, M_C2, F_C2, E_C2, R_C2, S_C2, H_C2, W_C2, U_C2);
143     bias(ofmap3, ofmap3, bias2, N_C2, M_C2, E_C2, F_C2);
144     pool(ofmap4, ofmap3, E_C2, F_C2, M_C2);
145
146     // Layer #3
147     //
148     convolution(ofmap5, ofmap4, fmap3, N_C3, C_C3, M_C3, F_C3, E_C3, R_C3, S_C3, H_C3, W_C3, U_C3);
149     bias(ofmap5, ofmap5, bias3, N_C3, M_C3, E_C3, F_C3);
150     relu(ofmap6, ofmap5, E_C3, F_C3, M_C3);
151
152     // Layer #4
153     //
154     convolution(ofmap, ofmap6, fmap4, N_C4, C_C4, M_C4, F_C4, E_C4, R_C4, S_C4, H_C4, W_C4, U_C4);
155     bias(ofmap, ofmap, bias4, N_C4, M_C4, E_C4, F_C4);
156
157     free(ofmap1);
158     free(ofmap2);
159     free(ofmap3);
160     free(ofmap4);
161     free(ofmap5);
162     free(ofmap6);
163
164     // To this line
165     // You are allowed to modify only the inside of this function
166     // From this line
167 }
168
169
170
171
```

Layers 1 & 2: All MACs in HW

Design Constraints

□ Example of accelerator

- Layers 2 & 3 in HW and the others in SW



* The FPGA considered here (ZC7020) is assumed to accommodate no more than **220 MAC hardware units** and **570KBs BRAM**

Evaluation

☐ Submission completeness

- Reproducibility (10pt)

☐ Accuracy

- Classification error (20pt)
- Quantization error (10pt)

☐ Performance

- Execution time (30pt)

☐ Technical solidness

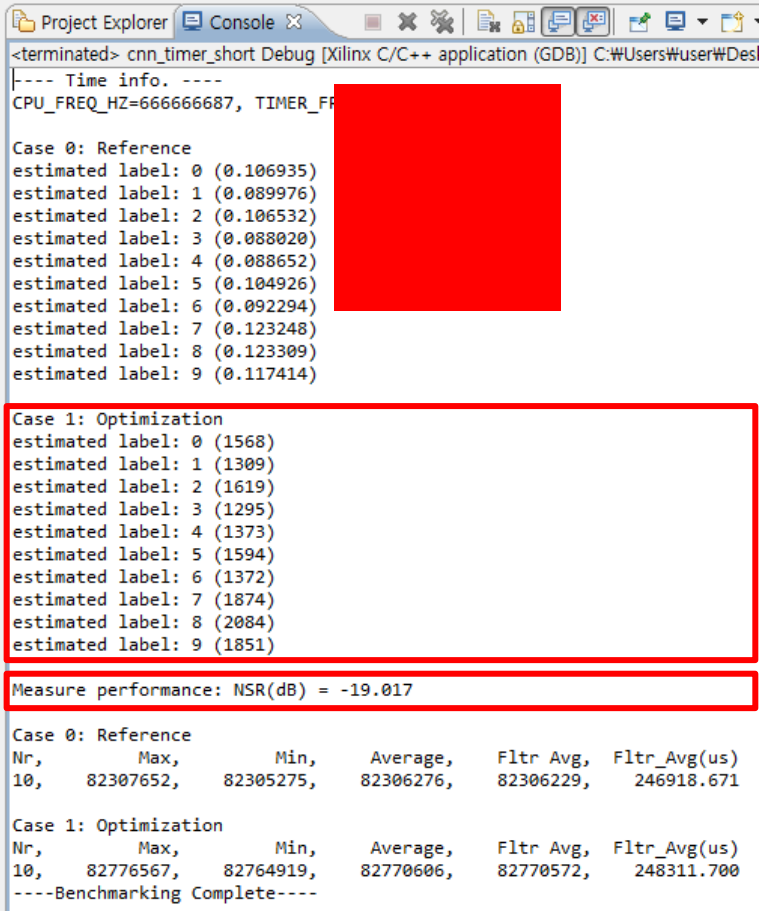
- Any good ideas or contributions (30pt)

Reproducibility

- ❑ Make sure that your submission is **reproducible**
 - In other words, it is possible to **reproduce** your design together with the claimed accuracy and performance using **only** the files that you submitted

Accuracy

- ❑ Run the main program to check the **accuracy** of **conv_hw**
 - Classification error
 - ✓ A **new** set of images will be given **onsite**
 - Noise-to-signal ratio (NSR)
 - ✓ Difference from **conv_ref**



The screenshot shows a console window with the following content:

```
<terminated> cnn_timer_short Debug [Xilinx C/C++ application (GDB)] C:\Users\User\Des
---- Time info. ----
CPU_FREQ_HZ=666666687, TIMER_F

Case 0: Reference
estimated label: 0 (0.106935)
estimated label: 1 (0.089976)
estimated label: 2 (0.106532)
estimated label: 3 (0.088020)
estimated label: 4 (0.088652)
estimated label: 5 (0.104926)
estimated label: 6 (0.092294)
estimated label: 7 (0.123248)
estimated label: 8 (0.123309)
estimated label: 9 (0.117414)

Case 1: Optimization
estimated label: 0 (1568)
estimated label: 1 (1309)
estimated label: 2 (1619)
estimated label: 3 (1295)
estimated label: 4 (1373)
estimated label: 5 (1594)
estimated label: 6 (1372)
estimated label: 7 (1874)
estimated label: 8 (2084)
estimated label: 9 (1851)

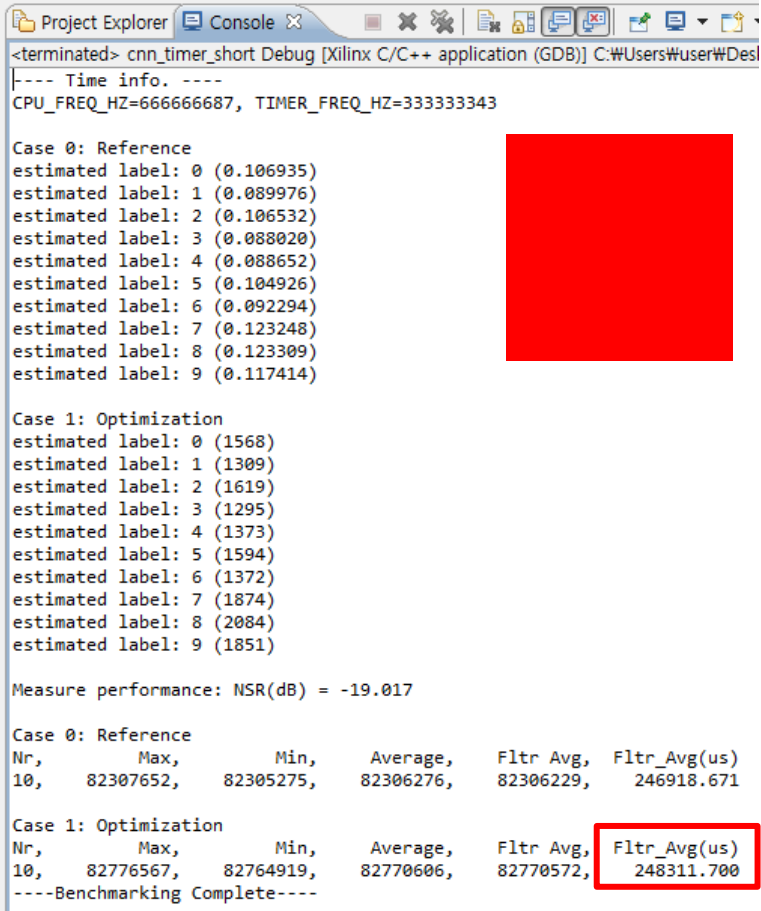
Measure performance: NSR(dB) = -19.017

Case 0: Reference
Nr,      Max,      Min,      Average,  Fltr Avg,  Fltr_Avg(us)
10,      82307652,    82305275,    82306276,    82306229,    246918.671

Case 1: Optimization
Nr,      Max,      Min,      Average,  Fltr Avg,  Fltr_Avg(us)
10,      82776567,    82764919,    82770606,    82770572,    248311.700
----Benchmarking Complete----
```


Performance

- ❑ Run the main program to check the **performance** of **conv_hw**
 - Execution time (-O3)



The screenshot shows a console window with the following output:

```
<terminated> cnn_timer_short Debug [Xilinx C/C++ application (GDB)] C:\Users\User\Des
|---- Time info. ----
CPU_FREQ_HZ=666666687, TIMER_FREQ_HZ=333333343

Case 0: Reference
estimated label: 0 (0.106935)
estimated label: 1 (0.089976)
estimated label: 2 (0.106532)
estimated label: 3 (0.088020)
estimated label: 4 (0.088652)
estimated label: 5 (0.104926)
estimated label: 6 (0.092294)
estimated label: 7 (0.123248)
estimated label: 8 (0.123309)
estimated label: 9 (0.117414)

Case 1: Optimization
estimated label: 0 (1568)
estimated label: 1 (1309)
estimated label: 2 (1619)
estimated label: 3 (1295)
estimated label: 4 (1373)
estimated label: 5 (1594)
estimated label: 6 (1372)
estimated label: 7 (1874)
estimated label: 8 (2084)
estimated label: 9 (1851)

Measure performance: NSR(dB) = -19.017

Case 0: Reference
Nr,      Max,      Min,      Average,  Fltr Avg,  Fltr_Avg(us)
10,      82307652,    82305275,    82306276,  82306229,  246918.671

Case 1: Optimization
Nr,      Max,      Min,      Average,  Fltr Avg,  Fltr_Avg(us)
10,      82776567,    82764919,    82770606,  82770572,  248311.700
----Benchmarking Complete----
```

A red square is overlaid on the right side of the console window. The value '248311.700' in the 'Fltr_Avg(us)' column for Case 1 is highlighted with a red box.

Technical Solidness

- ❑ Any good ideas or contributions that help to improve the performance such as
 - Optimum allocation of accelerator resources (MAC units & BRAMs) across layers
 - Maximum utilization of accelerator resources
- ❑ For the state-of-the-art hardware accelerator for CNN, refer to the following paper (and those citing it – available at <http://ieeexplore.ieee.org>):
 - “Maximizing CNN accelerator efficiency through resource partitioning”, Y. Shen et al., ISCA 2017

Important Notes


- ❑ You should keep the reference implementation (`conv_ref`) **unchanged**
- ❑ The compiler optimization level should be set to **-O3** when the performance is measured
- ❑ You will have some **bonus points** if you implement some of the paper ideas cited in the previous slide
- ❑ You may be able to provide **two** versions of the convolution function, e.g.,
 - 1st version with the **best** performance
 - 2nd version implementing paper idea(s)

Submission

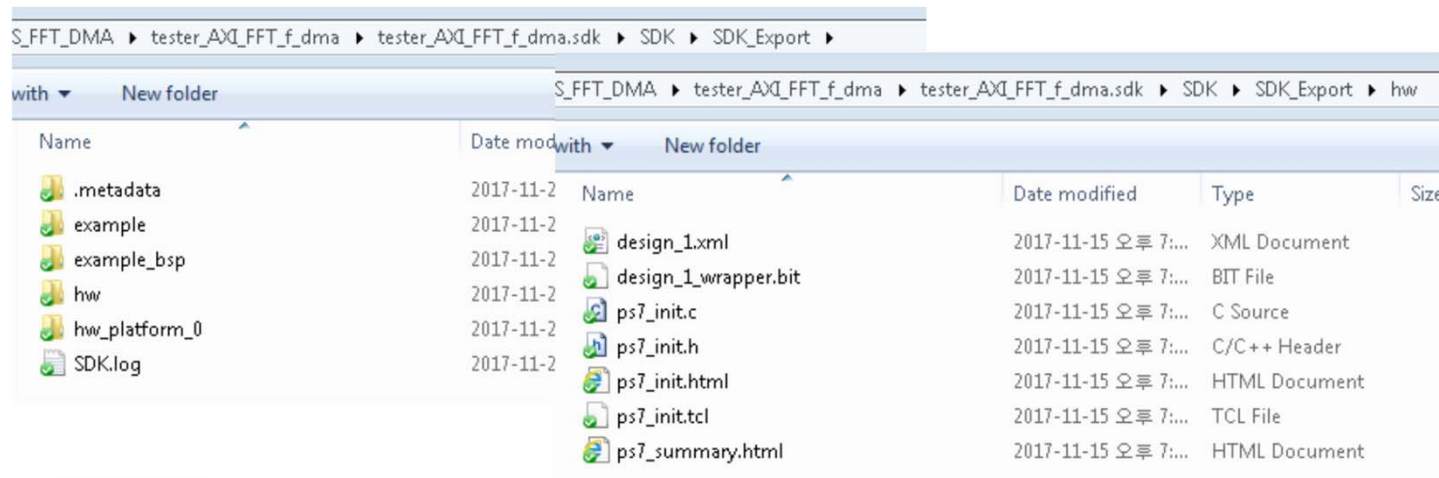
- ☐ Due date: **Dec. 15 (Sun), 2019, 15:00:00** GMT+9
- ☐ Zip the following files into a single file and send it to chesterku2013@gmail.com
 - **C/Verilog source/header files** and any other files that are needed to *reproduce* your design
 - Copy of the entire **SDK folder** in your project (including the bitstream)
 - **Slideset file** (PPT) including the results in the Console window
- ☐ You can post a question in the Q&A of the course page (<https://www.sites.google.com/site/kusocdesignlab/q-a-2>)
- ☐ **No delay** will be acceptable!

Submission

❑ SDK folder in your project



Note that the name of the SDK folder generally is given as “[project name].sdk”



Presentation & Demo

- ❑ **Dec. 16 (Mon) 10:30~13:30, New Eng. Bldg. #1113**
 - Team-presentation with exactly the same **slideset files** as submitted on **Dec. 15**
 - Team-demo with exactly the same **SDK folder** as submitted on **Dec. 15**
- 1. *Bring a storage device (e.g., HDD) having the entire project folder just in case (e.g., when the SDK folder does not work)*
- 2. *Note that any progress made later than **Dec. 15** **cannot** be counted for evaluation*

Appendix


Open Vivado projects

❑ Open Vivado projects

- Unzip and open '**TP_SD2019_v1**' folder and run file '**TP_SD2019_v1.xpr**'

Download link:

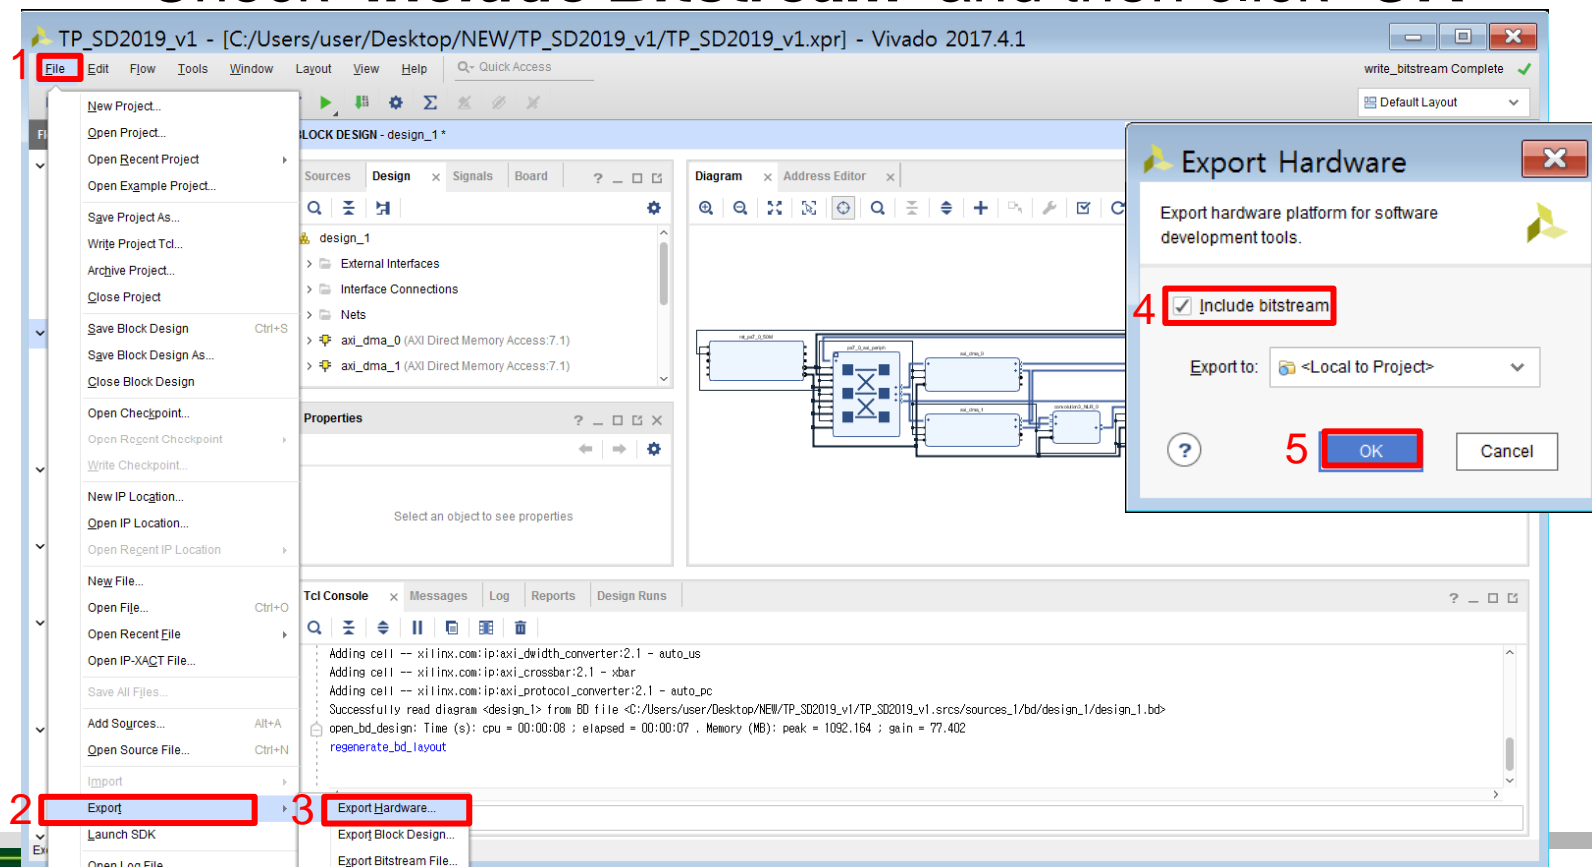
<https://drive.google.com/open?id=1zwLSdnKHZNEAO4TTdc0Yk2tIPeKX14SV>



Open Vivado projects

❑ Export Hardware

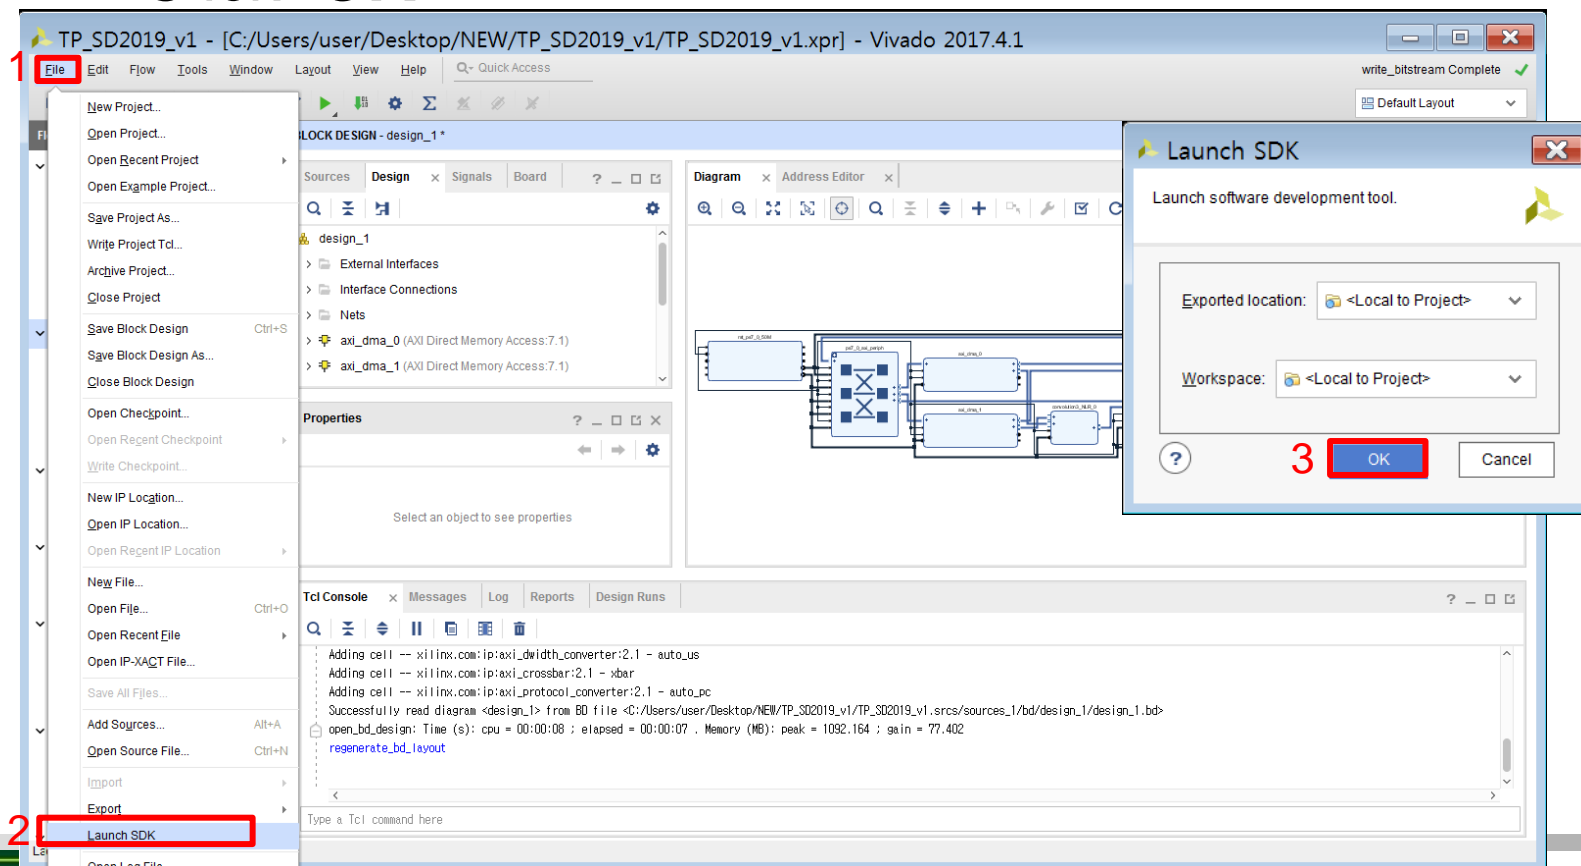
- Click '**Export > Export Hardware**' in '**File**' menu
- Check '**Include Bitstream**' and then click '**OK**'



Open Vivado projects

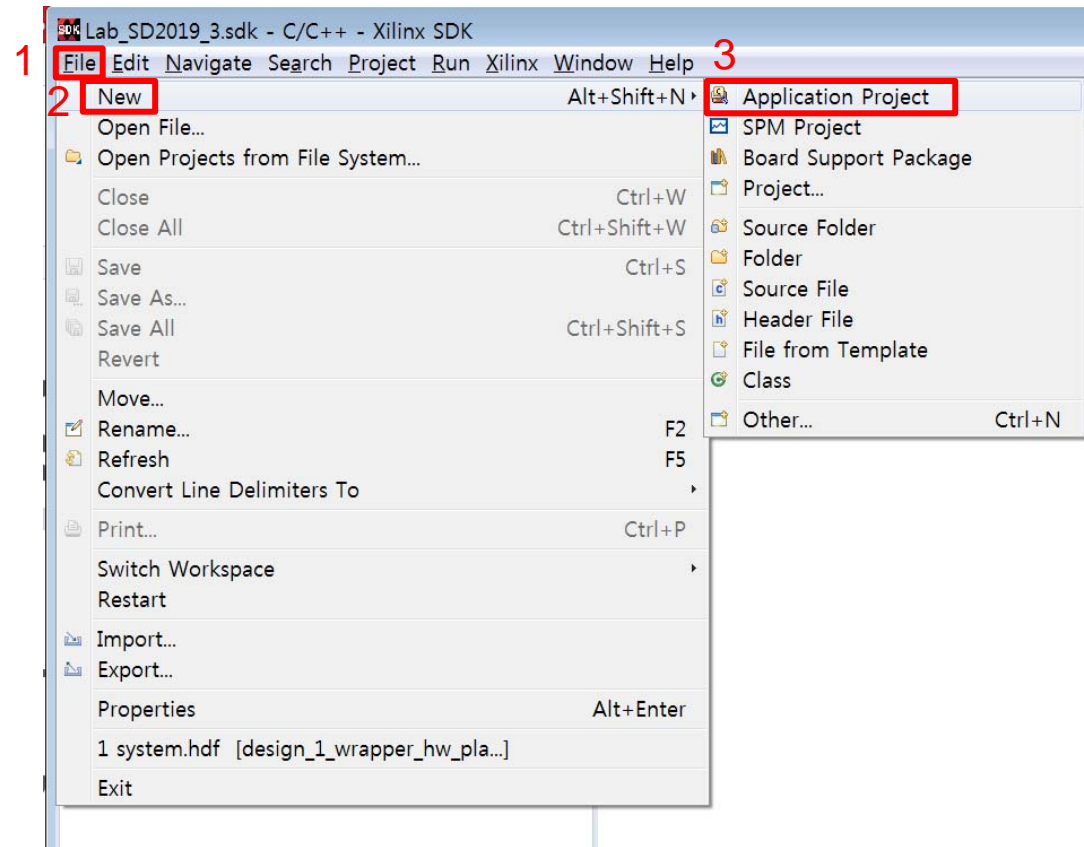
❑ Launch SDK

- Click '**Launch SDK**' in '**File**' menu
- Click '**OK**'



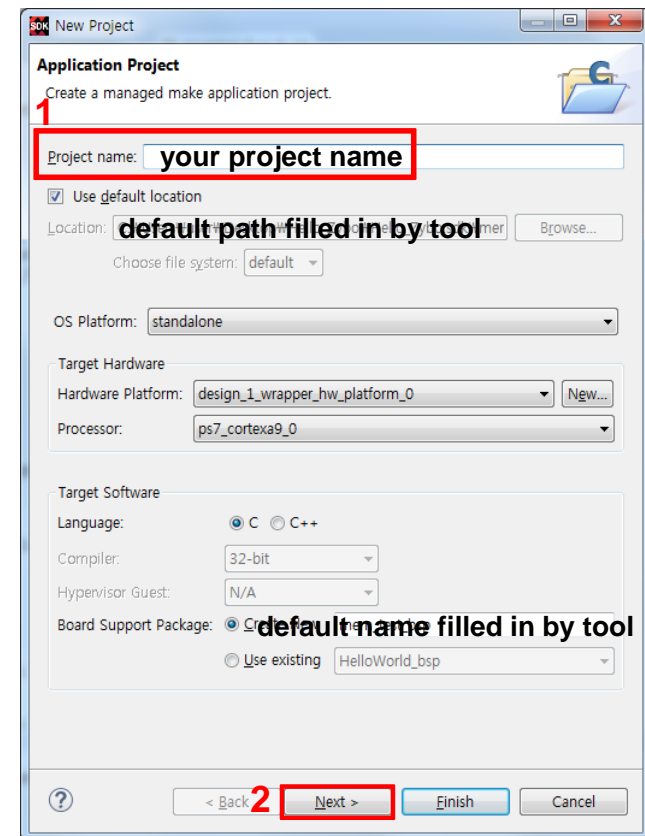
Running C Applications

- ❑ Create a C application project
 - Click '**File**' > '**New**' > '**Application Project**'




Running C Applications

- ❑ Create a C application project (cont'd)
 - Type the project name
 - Click '**Next**'



Running C Applications

- ❑ Create a C application project (cont'd)
 - Choose ***'Hello World'*** and then click ***'Finish'***



Running C Applications

❑ Add source files

- Choose 'your project name'
- Unzip and copy src folder and paste into the 'your project name' folder
- Click '**OK > Yes To All**'

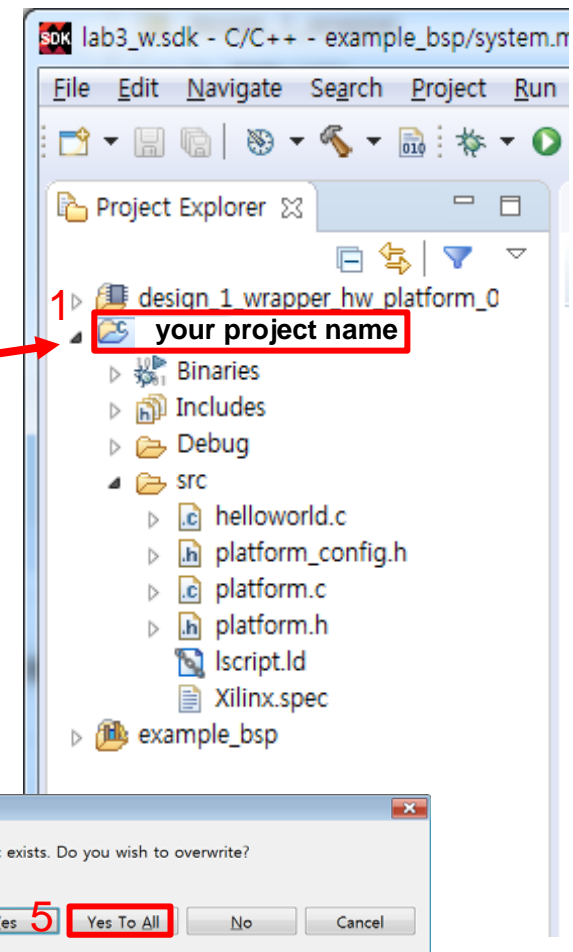


src.zip



src

2



3

4

1

your project name

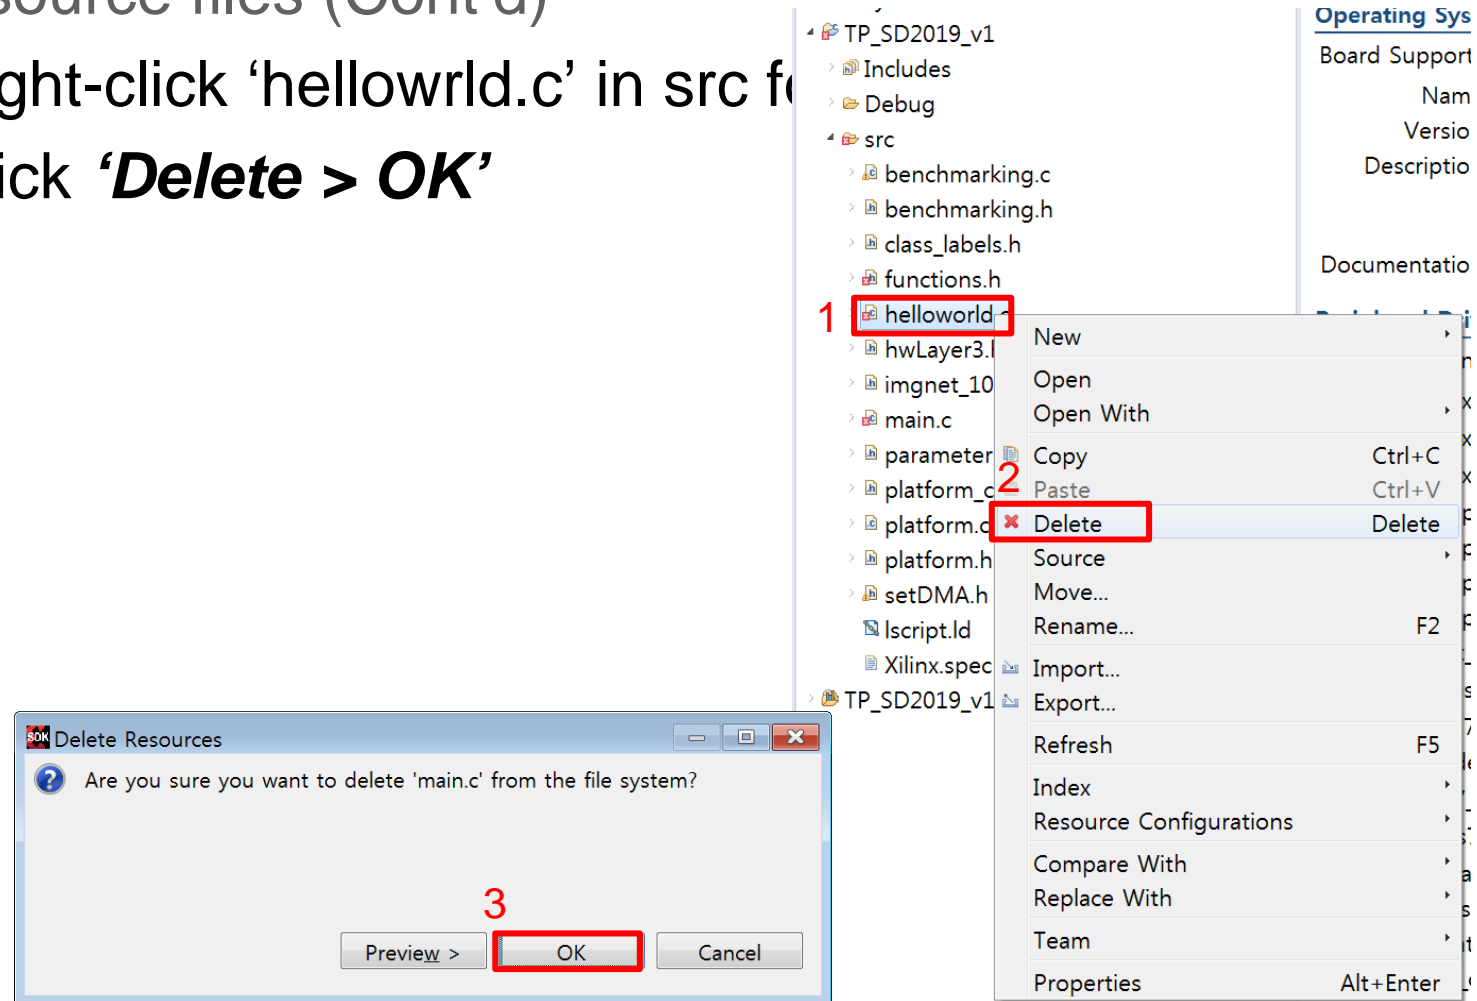
5

Yes To All

Running C Applications

❑ Add source files (Cont'd)

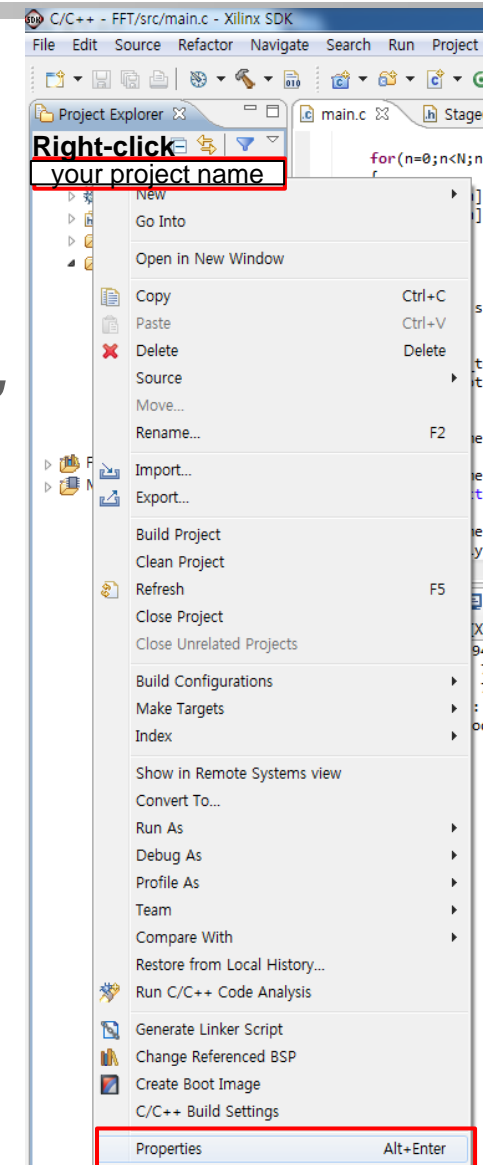
- Right-click 'helloworld.c' in src folder
- Click '**Delete > OK**'



Running C Applications

❑ Adding Math Library

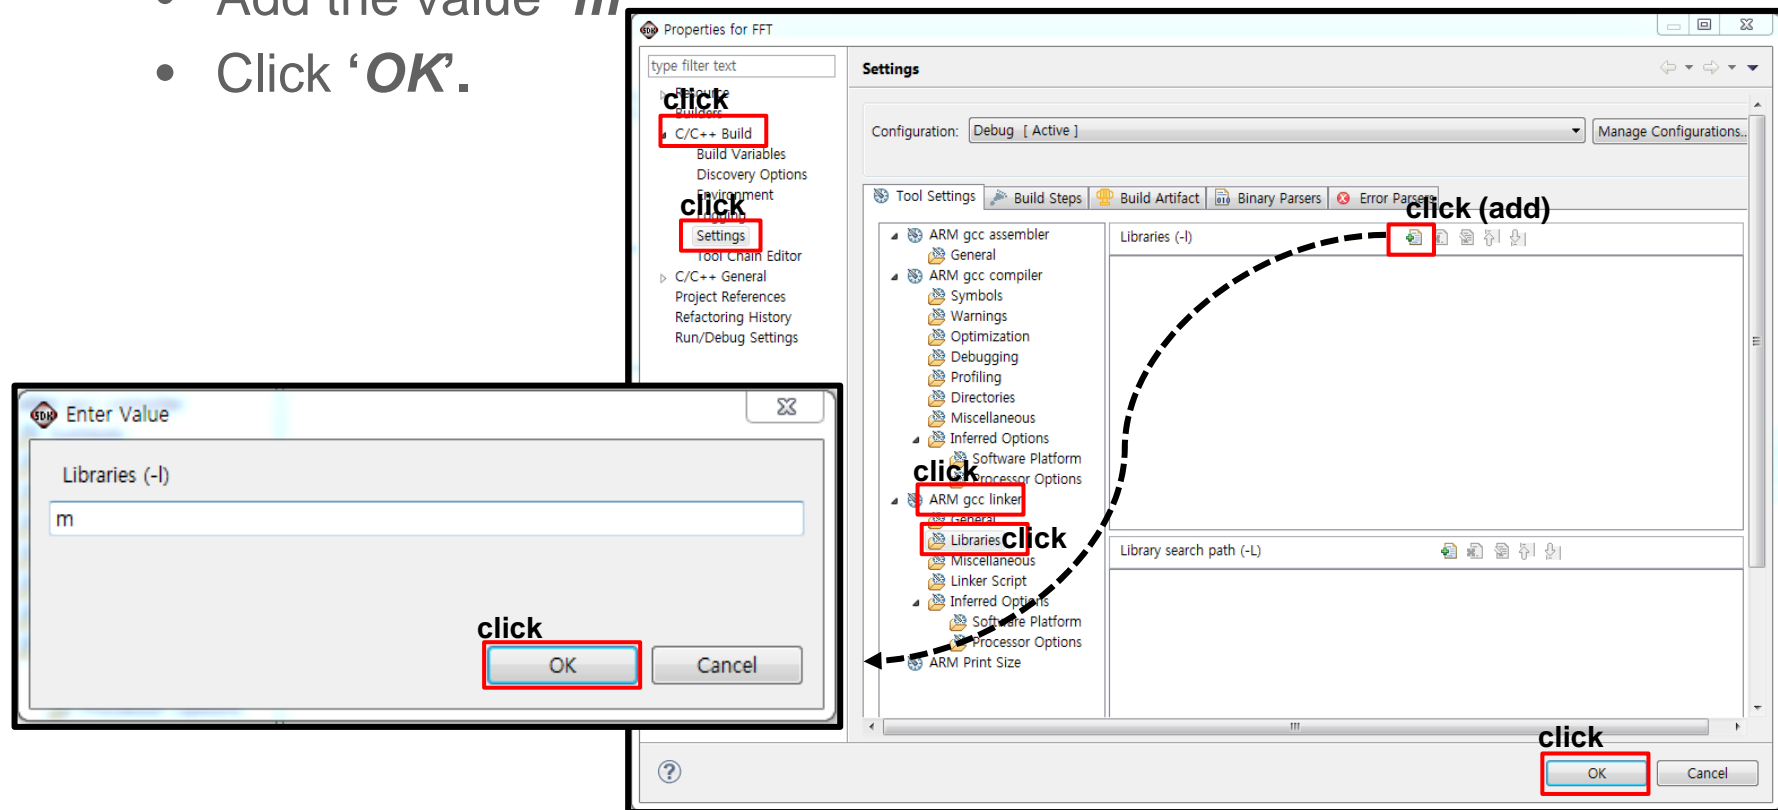
- Project must have *'-lm library'* to use *'math.h'* header file.
- Right-click *'your project name'* and then click *'Properties'*



Running C Applications

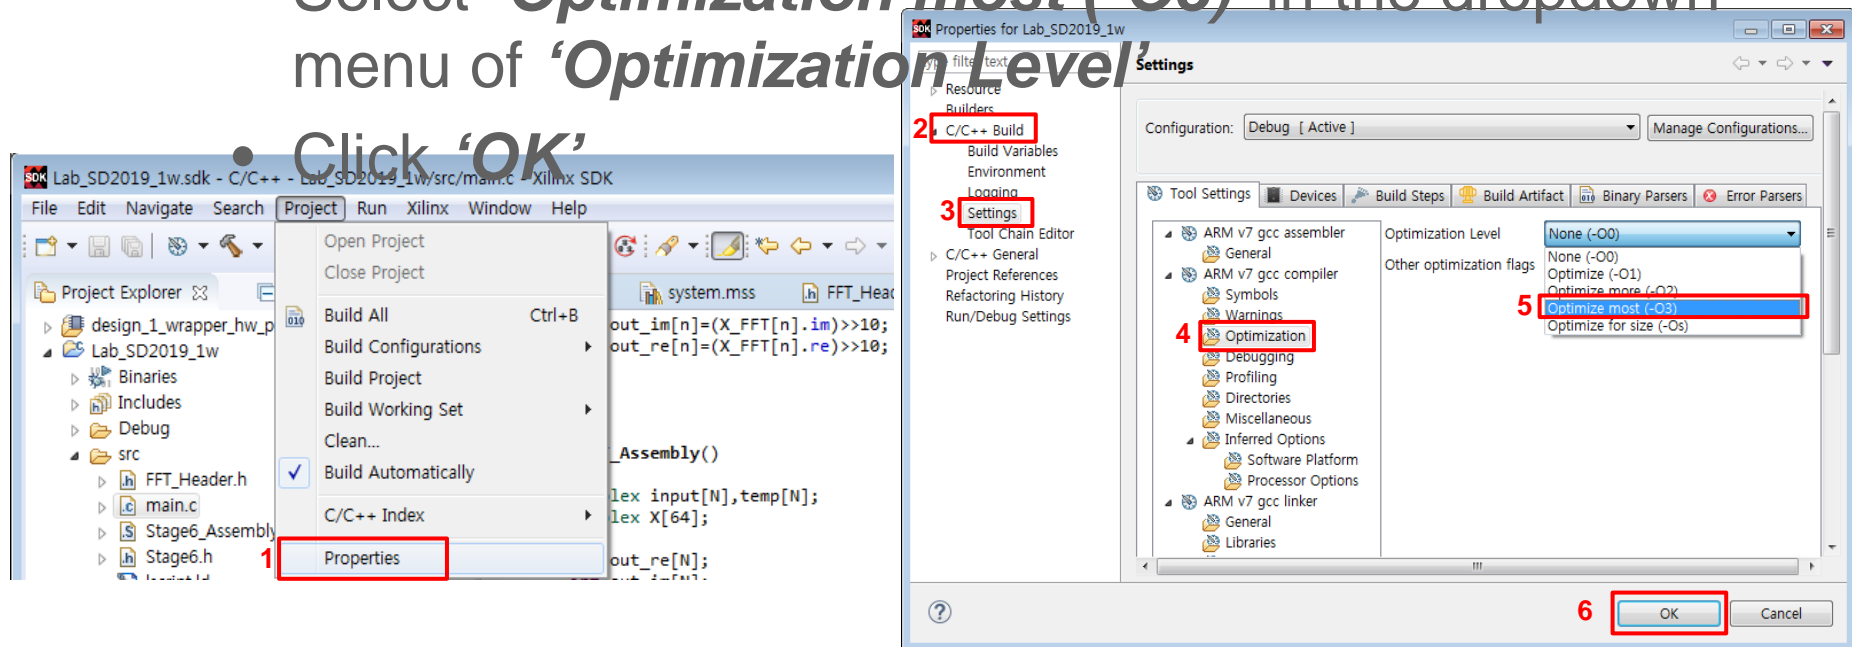
❑ Adding Math Library (Cont'd)

- Click '**C/C++ Build > Settings > ARM gcc linker > libraries > add**'
- Add the value '**m**'
- Click '**OK**'.




Running C Applications

- ❑ Set the compiler optimization level
 - Select '**Project**' menu and click '**Properties**'
 - Select '**Settings**' tap and click '**ARM v7 gcc compiler > Optimization**'
 - Select '**Optimization most (-O3)**' in the dropdown menu of '**Optimization Level**'
 - Click '**OK**'



Running C Applications

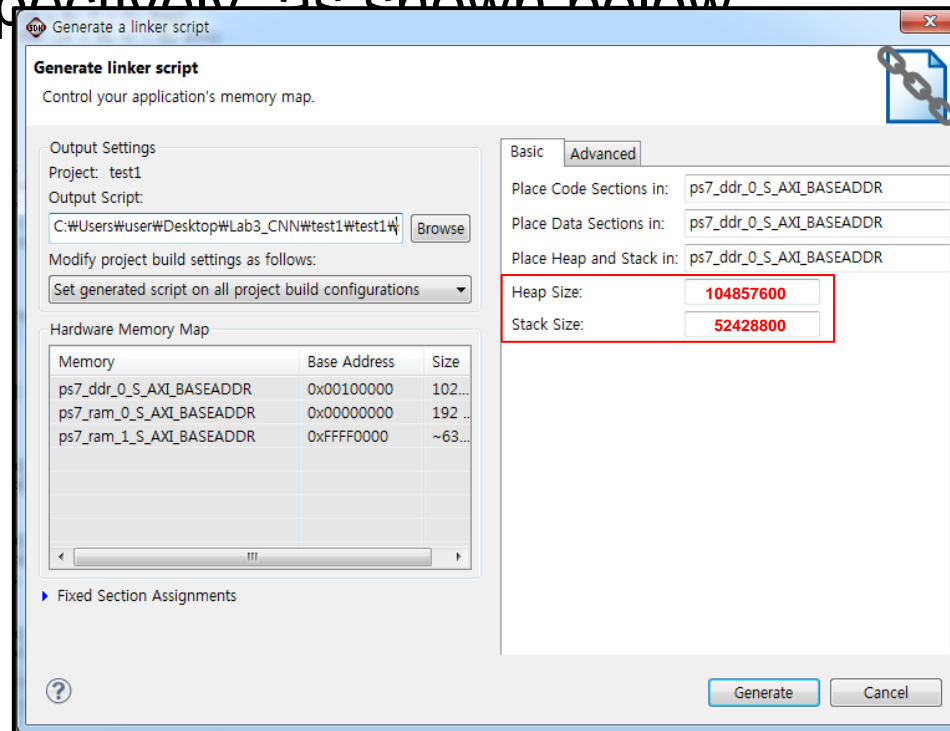
- ❑ Setting Stack & Heap Size
 - Select the application project in the **Project Explorer** or **C/C++ Projects** view
 - Right-click '**your project name**' and then click '**Generate Linker Script**' or click '**Xilinx Tools > Generate Linker script**'



Running C Applications

❑ Setting Stack & Heap Size (Cont'd)

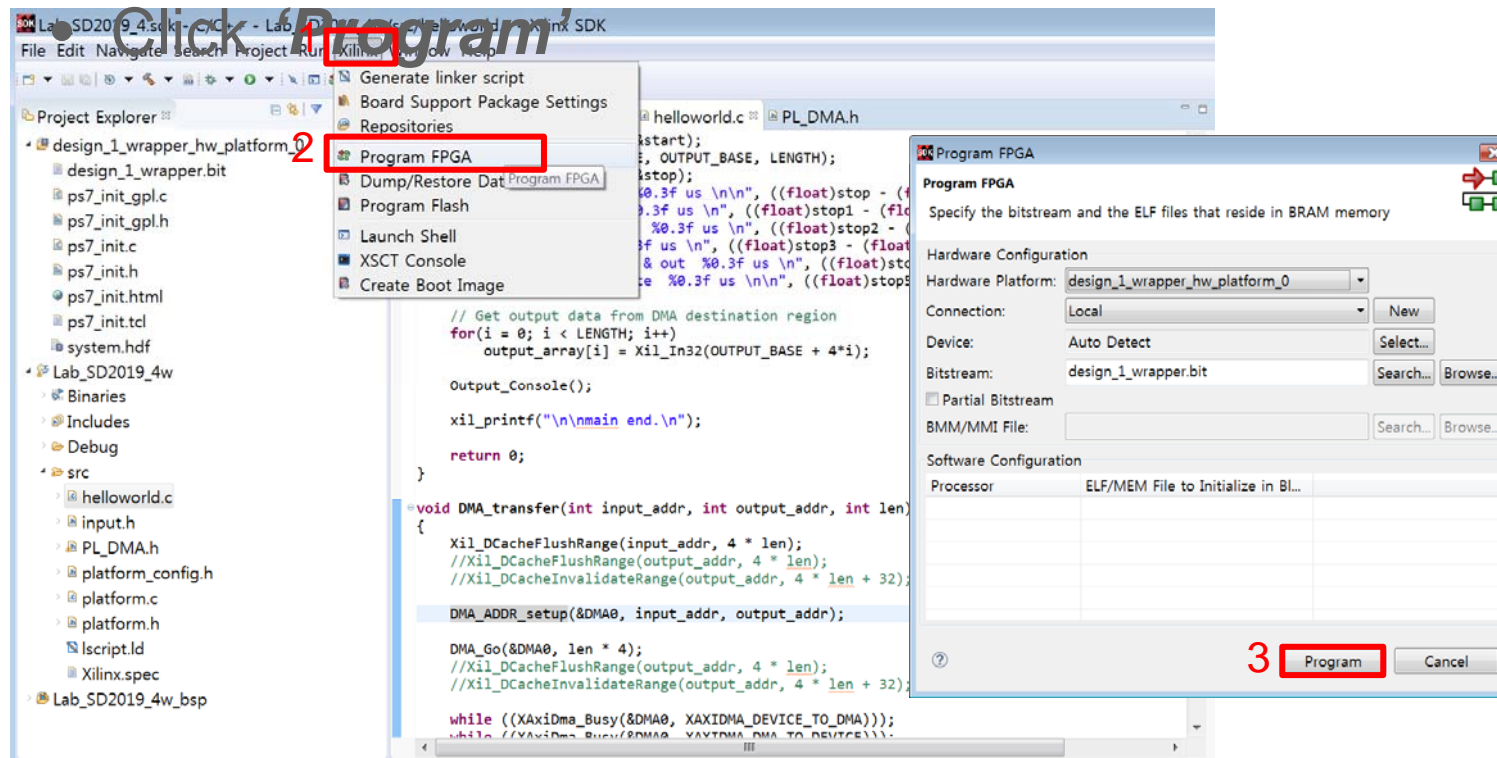
- Set both the heap and stack sizes in the **Basic** tab to **104857600** (100 MB) and **52428800** (50 MB) respectively, as shown below.



Running C Applications

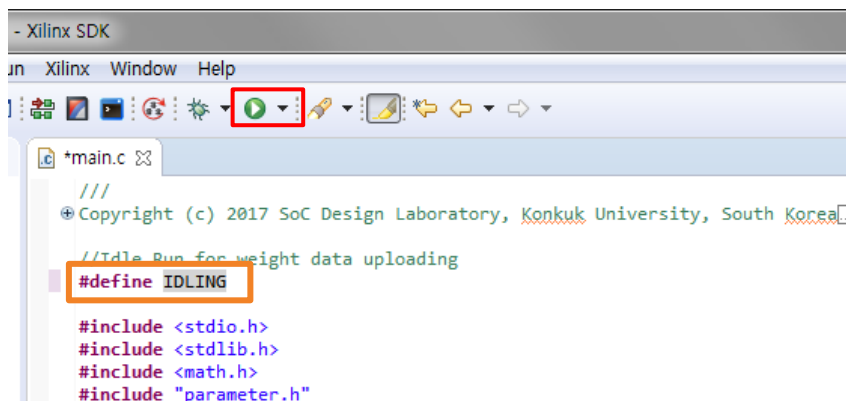
❑ Program FPGA

- Choose the '**Xilinx**' menu and then click '**Program FPGA**'



Running C Applications

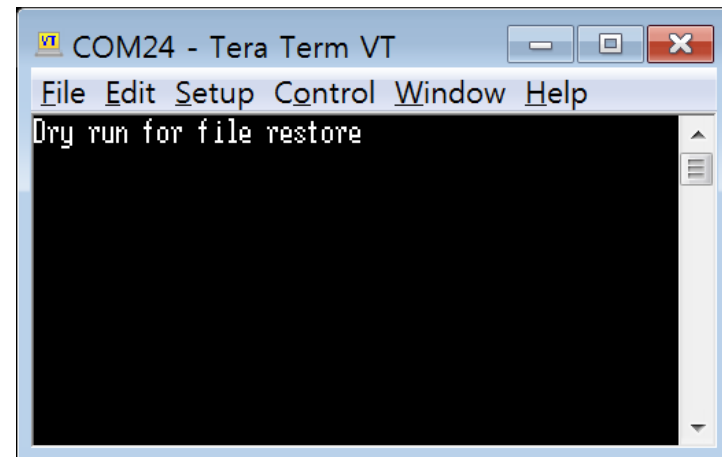
- ❑ Check the source code: **main.c**
 - Make sure that '**#define IDLING**' is written at the top of the source code
- ❑ Run the application
 - This application is simply for uploading the convolution weights into the DDR



```
///
Copyright (c) 2017 SoC Design Laboratory, Konkuk University, South Korea

//Idle Run for weight data uploading
#define IDLING

#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include "parameter.h"
```

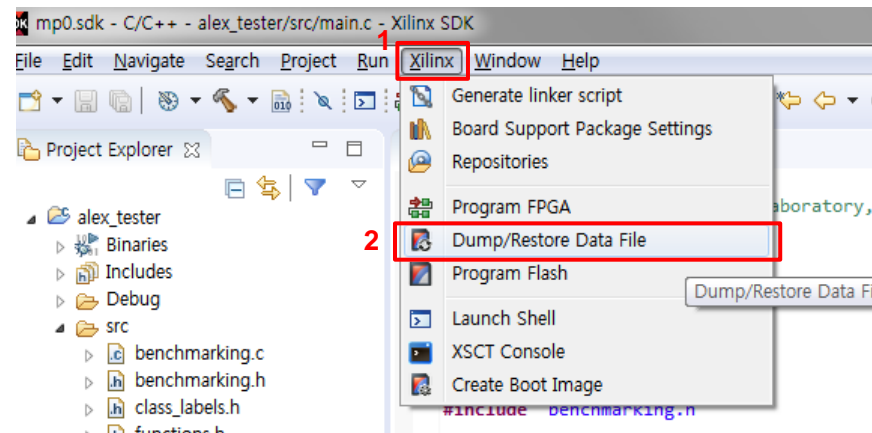


```
COM24 - Tera Term VT
File Edit Setup Control Window Help
Dry run for file restore
```

Running C Applications

❑ Restore Memory

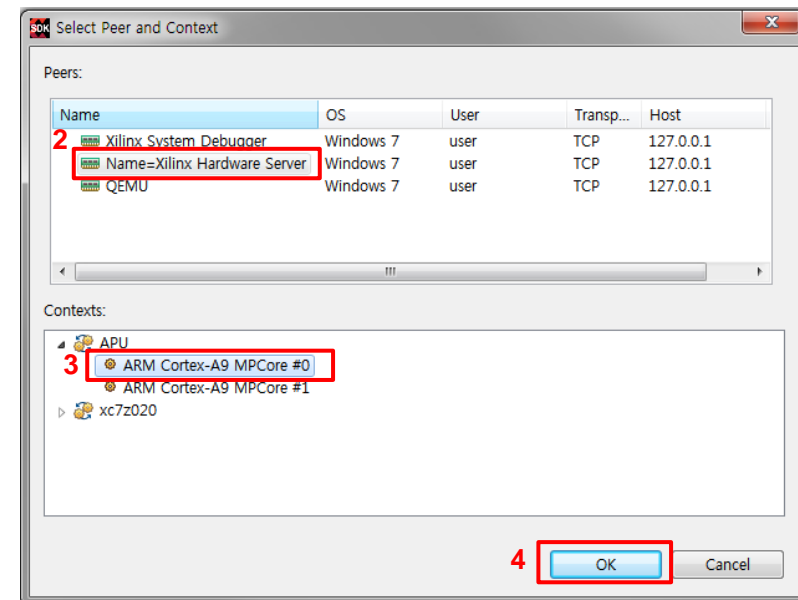
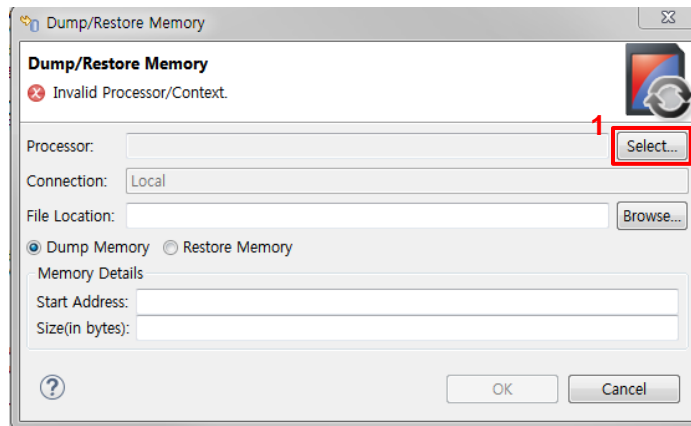
- Click '***Xilinx > Dump/Restore Data File***'



Running C Applications

❑ Restore Memory (Cont'd)

- Click '**Select**'
- Select '**Name=Xilinx Hardware Server**'
- Select '**ARM Cortex-A9 MPCore #0**'
- Click '**OK**'



Running C Applications

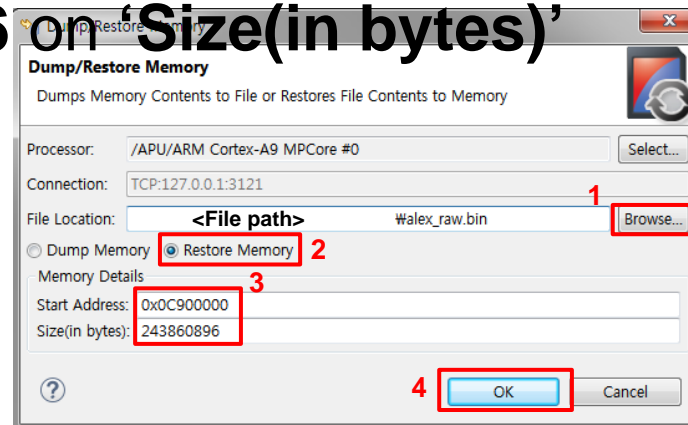
❑ Restore Memory (Cont'd)

- Click '**Browse**' and select file **alex_raw.bin**

Download link:

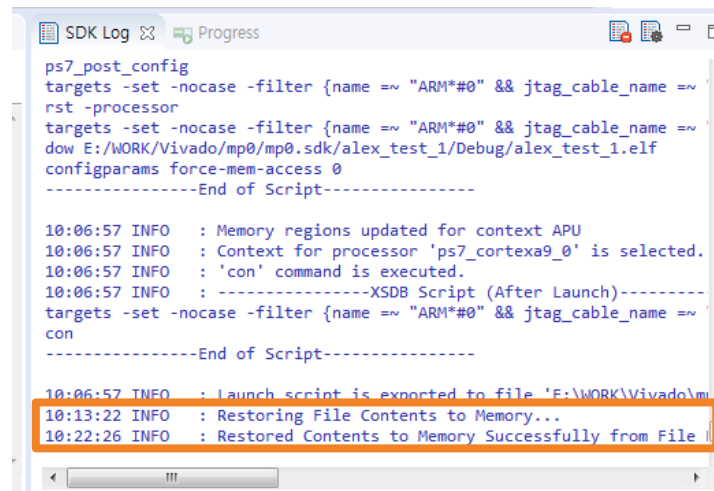
https://drive.google.com/open?id=122wv9t3FVrXlQYgc9bk7eK_L5-Q8ansA

- Select '**Restore Memory**'
- Type **0x0C900000** on 'Start Address' and **243860896** on 'Size(in bytes)'



Running C Applications

- ❑ Restore Memory (Cont'd)
 - Check the SDK log window (which may take about 9 minutes) to see the message '**Restored Contents to Memory Successfully from File ~**'



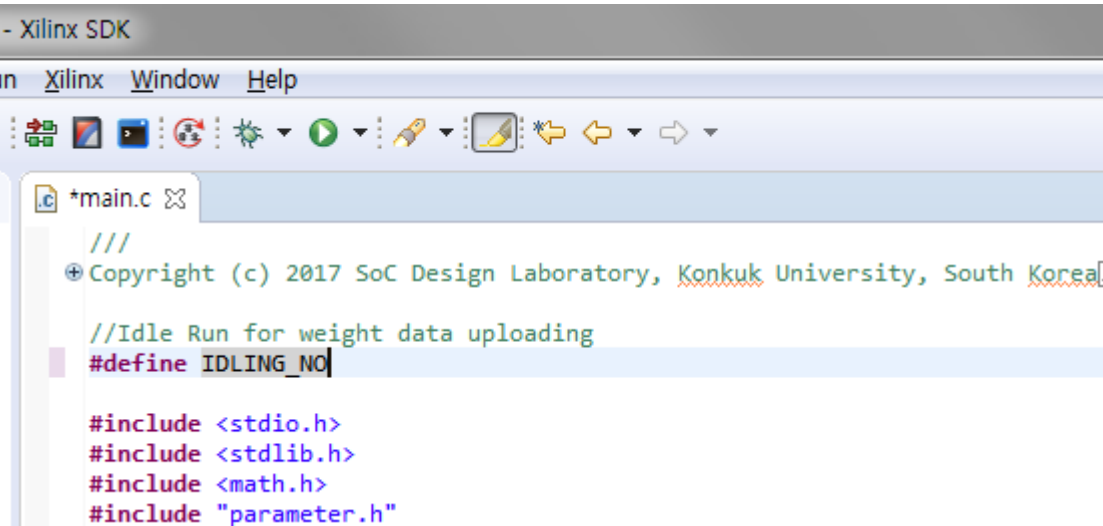
The screenshot shows the 'SDK Log' window with a 'Progress' tab. The log contains several lines of text, including configuration commands and status messages. A red rectangular box highlights the following three lines of log output:

```
10:06:57 INFO : Launch script is exported to file 'E:\WORK\Vivado\m
10:13:22 INFO : Restoring File Contents to Memory...
10:22:26 INFO : Restored Contents to Memory Successfully from File
```

Running C Applications

- ❑ Modify the source code: *main.c*
 - Change '**#define IDLING**' to '**#define IDLING_NO**' (or anything different from '**IDLING**')
- ❑ Run the application again
 - This application is for actually running the CNN

al

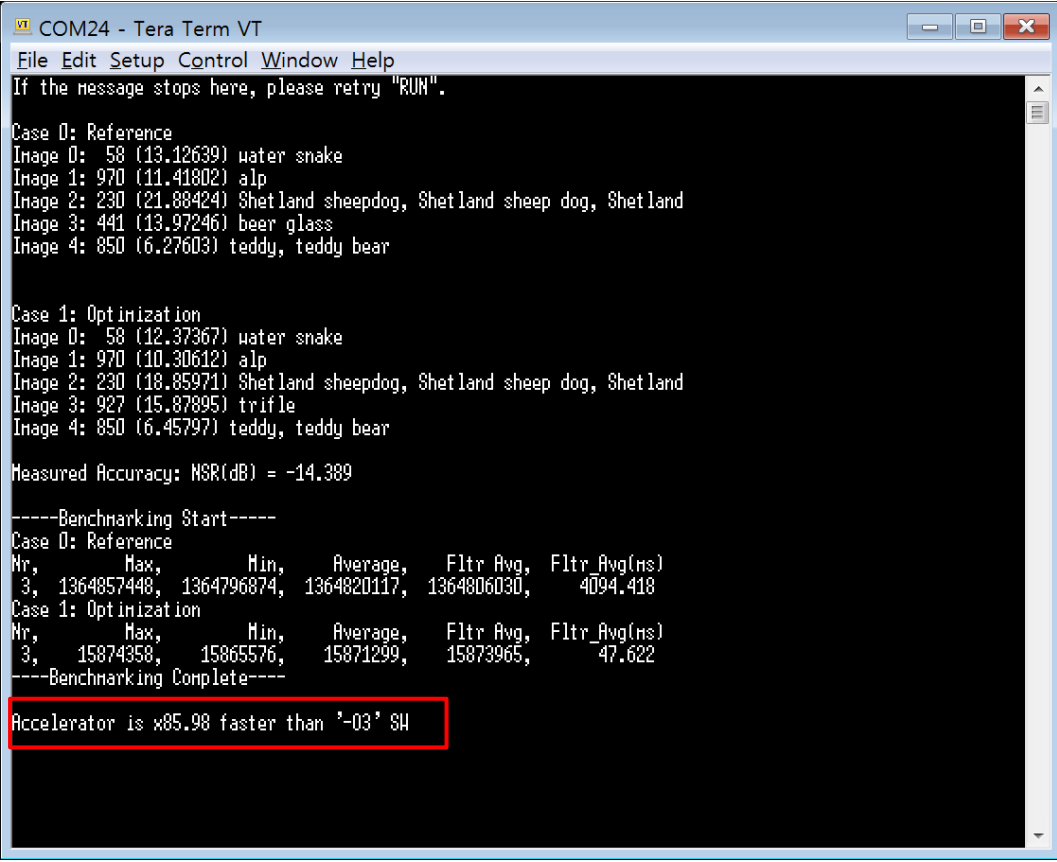


```
- Xilinx SDK
Xilinx Window Help
*main.c
///
+ Copyright (c) 2017 SoC Design Laboratory, Konkuk University, South Korea.
//Idle Run for weight data uploading
#define IDLING_NO

#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include "parameter.h"
```

Running C Applications

- ❑ Run the application again (Cont'd)
 - Check the output of the application on '*Tera Term*'



```
COM24 - Tera Term VT
File Edit Setup Control Window Help
If the message stops here, please retry "RUN".

Case 0: Reference
Image 0: 58 (13.12639) water snake
Image 1: 970 (11.41802) alp
Image 2: 230 (21.88424) Shetland sheepdog, Shetland sheep dog, Shetland
Image 3: 441 (13.97246) beer glass
Image 4: 850 (6.27603) teddy, teddy bear

Case 1: Optimization
Image 0: 58 (12.37367) water snake
Image 1: 970 (10.30612) alp
Image 2: 230 (18.85971) Shetland sheepdog, Shetland sheep dog, Shetland
Image 3: 927 (15.87895) trifle
Image 4: 850 (6.45797) teddy, teddy bear

Measured Accuracy: NSR(dB) = -14.389

-----Benchmarking Start-----
Case 0: Reference
Nr,      Max,      Min,      Average,  Fltr Avg,  Fltr Avg(ns)
3, 1364857448, 1364796874, 1364820117, 1364806030, 4094.418
Case 1: Optimization
Nr,      Max,      Min,      Average,  Fltr Avg,  Fltr Avg(ns)
3, 15874358, 15865576, 15871299, 15873965, 47.622
-----Benchmarking Complete-----

Accelerator is x85.98 faster than '-03' SW
```